



**UNIVERSITÀ  
DI TORINO**

**UNIVERSITÀ DEGLI STUDI DI TORINO**

Corso di Laurea Magistrale in Fisica dei Sistemi Complessi

# Improving VAE Representational Capabilities for Understanding LLMs

Tesi di Laurea Magistrale

**Relatore:**

Prof. Michele Caselle

**Corelatori:**

Prof. Alan Perotti

Prof. André Panisson

**Candidato:**

Francesco Marchisotti

**Tutor Aziendale:**

Gabriele Ciravegna

Anno Accademico 2025/2026

Francesco Marchisotti: *Improving VAE Representational Capabilities for Understanding LLMs*, Fisica dei Sistemi Complessi

**CENTAI TEAM:**

Gabriele Ciravegna

Alan Perotti

André Panisson

Francesco Cozzi

**UNITO PROFESSORS:**

Michele Caselle

Matteo Osella

**LOCATION:**

Torino

**TIME FRAME:**

Oct. 2025 - Jul. 2026

“I wish there was a way to know you’re in the good old days  
before you’ve actually left them.”

— *Andy Bernard*



# ABSTRACT

---

Concept-based methods represent a model’s computation in terms of human-meaningful concepts, but they typically encode each concept as a single scalar coefficient, e. g., a unit in the bottleneck of a Concept Bottleneck Model, or an atom in the dictionary of a Sparse AutoEncoder (SAE). This scalar encoding trades representational capacity for simplicity: each concept can vary only in magnitude, so reconstructing the underlying representation faithfully tends to require many concepts active at once. This thesis introduces the Variational Auto Embedding Encoder (VAEE), a model derived from first principles that represents each concept as a learned vector embedding instead, drawing together ideas from Concept Embedding Models, Variational AutoEncoders, and Vector-Quantized Variational AutoEncoders. The encoder maps an input to one vector embedding per concept; a learned prototype for each concept then decides, by similarity, which concepts are active, and the input is reconstructed from the stack of active embeddings. Its capabilities are demonstrated in one application, the decomposition of Large Language Model activations, where a VAEE with 8 active concepts matches the reconstruction of a parameter-matched TopK SAE that requires 128 — evidence that vector-valued concepts carry more information per unit of sparsity. The thesis is deliberate about the limits of this evidence: reconstruction was assessed only by mean squared error, each concept now spans several active neurons rather than one, and the recovered concepts were not found more interpretable than the SAE’s. Its contribution, then, lies in this added representational capacity; the interpretability of its concepts, their use for steering, and its application to representation learning more broadly are left for future work.



# ACKNOWLEDGMENTS

---

To whom shall be acknowledged,  
you are being acknowledged with this exact statement.

“I’m so funny.”

— *The candidate, who’s not actually funny*





# CONTENTS

---

|       |                                       |    |
|-------|---------------------------------------|----|
| 1     | INTRODUCTION                          | 1  |
| 1.1   | Motivation . . . . .                  | 1  |
| 1.2   | Thesis Structure . . . . .            | 1  |
| 2     | BACKGROUND AND RELATED WORK           | 3  |
| 2.1   | Natural Language Processing . . . . . | 3  |
| 2.1.1 | Language Modelling . . . . .          | 3  |
| 2.1.2 | The Transformer . . . . .             | 4  |
| 2.2   | eXplainable AI . . . . .              | 6  |
| 2.2.1 | Traditional XAI . . . . .             | 6  |
| 2.2.2 | Concept-based XAI . . . . .           | 8  |
| 2.3   | AutoEncoders . . . . .                | 10 |
| 2.3.1 | Variational AutoEncoders . . . . .    | 10 |
| 2.3.2 | Sparse AutoEncoders . . . . .         | 13 |
| 3     | METHODS                               | 19 |
| 3.1   | Aim of the Work . . . . .             | 19 |
| 3.2   | Notation . . . . .                    | 20 |
| 3.3   | Model Development . . . . .           | 21 |
| 3.3.1 | Probabilistic Formulation . . . . .   | 21 |
| 3.3.2 | Training Objective . . . . .          | 24 |
| 3.3.3 | Architectures . . . . .               | 26 |
| 4     | EXPERIMENTS                           | 27 |
| 4.1   | Experiment setup . . . . .            | 27 |
| 4.2   | Baselines . . . . .                   | 28 |
| 4.3   | Evaluation Metrics . . . . .          | 28 |
| 4.3.1 | Reconstruction . . . . .              | 29 |
| 4.3.2 | Sparsity . . . . .                    | 29 |
| 4.4   | Model Training . . . . .              | 30 |
| 4.5   | Results . . . . .                     | 31 |
| 4.5.1 | Quantitative results . . . . .        | 31 |
| 4.5.2 | Qualitative results . . . . .         | 32 |
| 4.5.3 | Ablation study . . . . .              | 32 |
| 5     | CONCLUSIONS                           | 35 |
| 5.1   | Contributions . . . . .               | 35 |
| 5.2   | Future Work . . . . .                 | 35 |
| A     | VAEE DERIVATIONS                      | 37 |
| A.1   | Optimal Gating Mechanism . . . . .    | 37 |
| A.2   | Explicit ELBO Derivation . . . . .    | 39 |
|       | BIBLIOGRAPHY                          | 43 |

## LIST OF FIGURES

---

|            |                                                                       |    |
|------------|-----------------------------------------------------------------------|----|
| Figure 2.1 | A decoder-only transformer drawn around the residual stream . . . . . | 5  |
| Figure 2.2 | Autoencoder and VAE latent spaces on MNIST                            | 11 |
| Figure 2.3 | A polysemantic neuron in InceptionV1 . . . . .                        | 14 |
| Figure 3.1 | The VAEE generative model . . . . .                                   | 22 |
| Figure 4.1 | VAEE vs SAE frontier: active concepts . . . . .                       | 31 |
| Figure 4.2 | VAEE vs SAE frontier: active neurons . . . . .                        | 32 |
| Figure 4.3 | Distribution of VAEE gate activations . . . . .                       | 33 |

## LIST OF TABLES

---

|           |                                                                             |    |
|-----------|-----------------------------------------------------------------------------|----|
| Table 4.1 | SetFit/ <a href="#">SST-2</a> dataset breakdown . . . . .                   | 28 |
| Table 4.2 | Sparsity metrics for <a href="#">SAE</a> and <a href="#">VAEE</a> . . . . . | 30 |

## LISTINGS

---

## ACRONYMS

---

|        |                                                  |                                                                                             |
|--------|--------------------------------------------------|---------------------------------------------------------------------------------------------|
| CB-LLM | Concept Bottleneck <a href="#">LLM</a> . . . . . | 9                                                                                           |
| CBM    | Concept Bottleneck Model . . . . .               | <a href="#">v</a> , <a href="#">8 f.</a> , <a href="#">18 f.</a>                            |
| CE     | Cross-Entropy . . . . .                          | <a href="#">3</a> , <a href="#">17</a> , <a href="#">29</a>                                 |
| CEM    | Concept Embedding Model . . . . .                | <a href="#">v</a> , <a href="#">8</a> , <a href="#">18 ff.</a>                              |
| ELBO   | Evidence Lower Bound . . . . .                   | <a href="#">12</a> , <a href="#">24 f.</a> , <a href="#">37</a> , <a href="#">39 f.</a>     |
| KL     | Kullback–Leibler . . . . .                       | <a href="#">12 f.</a> , <a href="#">24 ff.</a> , <a href="#">30</a> , <a href="#">40 f.</a> |

|                                                   |                              |
|---------------------------------------------------|------------------------------|
| LCBM Learnable Concept-based Model . . . . .      | 19 f.                        |
| LF-CBM Label-free <b>CBM</b> . . . . .            | 8 f.                         |
| LLM Large Language Model . . . . .                | v, 3 f., 7–10, 16, 19 f., 27 |
| MECHINT Mechanistic Interpretability . . . . .    | 10, 13 f., 16, 28            |
| MLP Multi-Layer Perceptron . . . . .              | 4 ff., 16 f.                 |
| MSE mean squared error . . . . .                  | v, 10, 12, 24, 29, 31        |
| SAE Sparse AutoEncoder . . . . .                  | v, 10, 13–20, 27–31          |
| SST-2 Stanford Sentiment Treebank . . . . .       | 27                           |
| VAE Variational AutoEncoder . . . . .             | v, 10–13, 18, 20, 24         |
| VAEE Variational Auto Embedding Encoder . . .     | v, 20 ff., 24, 27–32         |
| VQ-VAE Vector-Quantized Variational AutoEncoder . | v, 13, 18, 20                |



# INTRODUCTION

---

## 1.1 MOTIVATION

## 1.2 THESIS STRUCTURE



## BACKGROUND AND RELATED WORK

---

This chapter establishes the background on which the thesis builds and reviews the related work, moving from the setting in which the contribution is demonstrated to the ideas from which it is built. Section 2.1 introduces Large Language Models and the transformer architecture that produces the dense internal activations on which that demonstration rests. Section 2.2 reviews explainable AI, tracing the field’s shift from attributing a prediction to input features towards explaining it in terms of human-meaningful *concepts*, the view this work adopts. Section 2.3 then develops the autoencoder family, from the standard autoencoder through the variational and sparse variants between which the thesis’s contribution sits. A single tension runs through all three: reconstructing a model’s activations faithfully while decomposing them into a small, interpretable set of concepts — the tension the method of section 3.3 sets out to relax.

### 2.1 NATURAL LANGUAGE PROCESSING

Natural Language Processing is the branch of machine learning concerned with getting computers to process and generate human language. For most of its history the field advanced through task-specific systems — separate models for translation, sentiment analysis, or question answering — but it has since converged on a single recipe: train one large neural network to predict text on a vast corpus, then reuse it across tasks. The Large Language Models (LLMs) that result are the systems on which the thesis’s contribution is demonstrated. What follows introduces only as much of their training and architecture as the later chapters rely on, with the emphasis on the internal activations these models compute rather than the text they emit.

#### 2.1.1 LANGUAGE MODELLING

A language model defines a probability distribution over sequences of text. In the dominant *autoregressive* formulation, this distribution is factorised left to right, so that the model repeatedly predicts the next unit of text given all the units before it [30]. The units are *tokens*, words or sub-word fragments drawn from a fixed vocabulary, and the model is a neural network that maps a sequence of tokens to a probability distribution over the vocabulary for the next position. Training minimises the Cross-Entropy (CE) between this predicted

distribution and the true next token, averaged over a large corpus; the objective requires no human annotation, since the supervision signal is the text itself.

Scaling this recipe to more parameters, more data, and more computing power yields LLMs whose capabilities extend well beyond the literal next-token task they were trained on. A network trained only to continue text acquires the ability to translate, summarise, answer questions, and follow instructions, often from nothing more than a description and a few examples supplied in its input, with no change to its weights [5, 7]. This generality is what makes LLMs both powerful and hard to interpret: a single fixed network comes to encode a vast range of linguistic and world knowledge, distributed across its parameters and activations in a form that no direct reading of the weights makes legible.

### 2.1.2 THE TRANSFORMER

Modern LLMs are almost universally built on the *transformer* architecture [44]. It embeds the input tokens into a sequence of vectors, one per position, and refines them through a stack of identical *blocks* before reading the final vectors out as next-token distributions. Each block is made of two sublayers with complementary roles: an *attention* sublayer, which exchanges information between positions, and a Multi-Layer Perceptron (MLP) sublayer, which transforms each position on its own. Figure 2.1 shows this organisation.

**ATTENTION** *Attention* is the mechanism by which positions exchange information. At each position, an *attention head* forms a query and compares it against a key emitted by itself and every position before it; the better a position matches, the more it contributes to a weighted average of the corresponding values, which becomes the head’s output for the current position [44]. Concretely, the head projects each stream vector to a query, key, and value,  $q_t = Qx_t$ ,  $k_s = Kx_s$ , and  $v_s = Vx_s$ , and returns

$$\text{head}(x_{1:t}) = \sum_{s=1}^t a_{ts} v_s, \quad a_{ts} = \text{softmax}_s \left( \frac{q_t^\top k_s}{\sqrt{d_k}} \right), \quad (2.1)$$

where  $d_k$  is the dimension of the queries and keys, and the softmax runs over  $s = 1, \dots, t$ , so that a position attends only to itself and those before it. A sublayer runs several such heads in parallel — each free to specialise, one linking a pronoun to its referent, another a verb to its subject — and combines their outputs by concatenation and a learned output projection. Because attention is the only component that moves information between positions, it is where the model assembles context: it turns a position’s representation from a fact about a single token into one that reflects the surrounding text.






Figure 2.1: A decoder-only transformer, drawn around the *residual stream*. Each token is embedded into a vector; the resulting per-position state flows up the central stream, and every block reads from the stream and adds its result back through a residual connection — first a multi-head self-attention sublayer, which mixes information across positions, then a position-wise MLP sublayer. After  $N$  such blocks the final stream state is unembedded into a next-token distribution. The additive structure of eq. (2.3) is what makes the residual stream a single, well-defined activation space to decompose; the view follows Elhage et al. [9].

**MLP SUBLAYERS** Where attention mixes information across positions, an MLP sublayer processes each position on its own, applying the same feed-forward transformation to the vector at every position independently,

$$\text{MLP}(\mathbf{x}_t) = W_2 \varphi(W_1 \mathbf{x}_t + \mathbf{b}_1) + \mathbf{b}_2, \quad (2.2)$$

with a pointwise nonlinearity  $\varphi$ , an up-projection  $W_1$  to a wider hidden layer and a down-projection  $W_2$  back. Much of the model’s capacity sits in these sublayers, which are widely thought to store and retrieve its learned associations and are a primary site at which interpretable features appear [4].

**THE RESIDUAL STREAM** A single detail ties these sublayers together. Neither sublayer overwrites the vector it is given; each computes an update and *adds* it back, through a *residual connection*. Every position therefore carries one running vector — the *residual stream* — that every sublayer of every block reads from and writes to. Let  $\mathbf{x}_t^{(\ell)}$  denote the

residual stream at position  $t$  as it enters block  $\ell$ ; the block updates it as

$$\mathbf{x}_t^{(\ell+1)} = \mathbf{x}_t^{(\ell)} + \text{Attn}^{(\ell)}(\mathbf{x}_{1:t}^{(\ell)}) + \text{MLP}^{(\ell)}(\mathbf{x}_t^{(\ell)}), \quad (2.3)$$

where  $\mathbf{x}_{1:t}^{(\ell)} = (\mathbf{x}_1^{(\ell)}, \dots, \mathbf{x}_t^{(\ell)})$  is the prefix of stream states up to position  $t$  — so the attention term, unlike the MLP term, depends on the whole prefix rather than on  $\mathbf{x}_t^{(\ell)}$  alone. The stream starts as the token’s embedding and accumulates the output of every sublayer in turn; because attention writes into it, a position’s stream is never a fact about that token alone but a running summary of the context around it. This additive structure is what makes the residual stream a single, well-defined space in which to study the model’s computation — the view developed in the mechanistic-interpretability literature [9] and adopted throughout this thesis. It is also the object the later chapters take apart: the activation vector  $\mathbf{x} \in \mathbb{R}^n$  that the autoencoders of section 2.3 receive as input is exactly such a residual stream state  $\mathbf{x}^{(\ell)}$ , read off at a chosen layer  $\ell$ .

## 2.2 EXPLAINABLE AI

The incredible performance of modern machine-learning systems comes at the cost of transparency: a deep network’s output is the product of millions of interacting parameters, with no accompanying account of why it was produced. eXplainable Artificial Intelligence (XAI) is the field that sets out to supply such accounts, making a model’s behaviour intelligible enough to be trusted, debugged, or learned from. The methods relevant to this thesis are most naturally organised around a shift in what counts as an explanation: from *attribution*, which asks *where* (“which input features drove a prediction?”), to *concepts*, the human-meaningful units in terms of which a model can be said to reason. It is this latter, concept-based view on which this work builds.

### 2.2.1 TRADITIONAL XAI

#### 2.2.1.1 Intrinsically Interpretable Models

The most direct route to an interpretable model is to choose one whose decision process is transparent by construction. Linear and logistic regression, decision trees, naive Bayes classifiers, and  $k$ -nearest neighbours are the canonical examples: their parameters or structure can be inspected directly, so that the explanation *is* the model rather than a separate artefact derived from it [26]. Rudin [35] argues that such models should be preferred wherever they are competitive, since post-hoc explanations of an opaque model are approximations that need not be faithful to its actual computation. This route is, however,

closed for the systems studied in this thesis: LLMs comprise billions of nonlinearly interacting parameters, and no direct reading of their weights yields a human-understandable account of their behaviour. Interpretability must therefore be recovered *post hoc*, from a model that is already trained and fixed.

#### 2.2.1.2 *Post-hoc Attribution Methods*

The dominant post-hoc paradigm explains an individual prediction by *attributing* it to the input features, assigning each feature a score that quantifies its contribution to the output. LIME [34] does this by fitting a simple, interpretable surrogate (typically a sparse linear model) to the behaviour of the black box in a local neighbourhood of the instance being explained. SHAP [23] unifies this and several related schemes under the Shapley value from cooperative game theory, yielding feature attributions that are additive and satisfy a set of desirable axioms. For models operating on images, gradient-based saliency maps [38] and their refinements such as Grad-CAM [36] produce a heatmap over the input that highlights the regions most influential to the prediction. What these methods share is also their central limitation: they indicate *where* in the input the model is sensitive, but not *what* the model has internally come to represent. An attribution map can mark a pixel or a token as important without revealing the human-meaningful concept it participates in: a gap that motivates the concept-based methods of section 2.2.2.

**COUNTERFACTUAL EXPLANATIONS** A counterfactual explanation answers a contrastive question: what is the smallest change to the input that would have produced a different, desired prediction [45]? Rather than scoring the features of the present input, it points to a nearby instance on the other side of the decision boundary: “had this applicant’s income been slightly higher, the loan would have been granted”. Such explanations are attractive because they are actionable and require no access to the model’s internals, only the ability to query it. Generating counterfactuals that are at the same time close to the original input and realistic is itself a modelling problem, one which generative models are naturally suited to solve.

#### 2.2.1.3 *Prototypes*

A complementary, example-based strategy explains a model not through feature scores but through representative instances. A *prototype* is a data point that is typical of a region of the input space, chosen so that a small set of such points summarises the data (or the model’s behaviour upon it) in terms a human can grasp by inspection [18]. The defining characteristic, for the purposes of this thesis, is that

a prototype in this traditional sense is an *input instance*: an actual or representative example drawn from the data space.

### 2.2.2 CONCEPT-BASED XAI

Concept-based XAI answers the *what* question that attribution leaves open: instead of pointing to the inputs responsible for a prediction, it explains the prediction in terms of high-level, human-meaningful *concepts* (“striped”, “has wheels”, “expresses uncertainty”) of the kind a person would use to describe the same decision. Methods in this family differ chiefly in whether the concepts are built into the model from the outset or recovered from a model that was trained without them.

#### 2.2.2.1 Explainable-by-design Models

One way to guarantee that concepts genuinely take part in a model’s decision is to make them an explicit, intermediate stage of the computation, so that the prediction is forced to pass through a concept representation before reaching the output.

**CONCEPT BOTTLENECK MODELS** Concept Bottleneck Models (CBMs) [22] insert exactly such a stage: the network first maps the input to a set of human-specified concepts and then predicts the label from those concepts alone. Because the output depends on the input only through the concept layer, the model’s reasoning can be read off the activated concepts, and a user can intervene on them, correcting a mistaken concept and observing the effect on the prediction. The downside is supervision: training a CBM requires data annotated by human experts with concept labels, which is costly and ties the model to a fixed, predefined vocabulary.

**CONCEPT EMBEDDING MODELS** A known weakness of CBMs is an accuracy–interpretability trade-off: when the predefined concepts do not carry all the information the task requires, forcing the prediction through a narrow concept bottleneck degrades performance. Concept Embedding Models (CEMs) [11] loosen the bottleneck by representing each concept not as a single scalar but as a pair of high-dimensional embeddings, one for its active and one for its inactive state. This recovers accuracy comparable to an unconstrained black box while retaining concept-level interpretability and the possibility of intervention. Like CBMs, these models require a concept-annotated dataset.

**LABEL-FREE CONCEPT BOTTLENECK MODELS** Label-free CBMs (LF-CBMs) [27] address the annotation cost directly. Rather than relying on a hand-labelled concept set, they extract candidate concepts from an LLM and align the network’s activations with them through

a vision–language model such as CLIP, converting a standard backbone into a CBM without any concept supervision. This makes the approach scalable to settings where dense concept annotation would be impractical.

**CONCEPT BOTTLENECK LLMs** The concept-bottleneck idea has recently been carried over to language. Concept Bottleneck LLMs (CB-LLMs) [40] insert a concept bottleneck layer into an LLM, routing text classification — and, in a generative variant, generation itself — through a set of interpretable concepts before the output. As in an LF-CBM, a hand-annotated concept set is not needed: the concepts are generated by a prompted LLM and each training text is scored against them automatically, so the bottleneck is built without concept supervision. The explicit concept layer also makes the model *steerable*, since intervening on a concept predictably shifts its behaviour. Like the other methods in this group, CB-LLMs are *explainable by design*: interpretability is secured by training the model with the concept bottleneck built in.

Even so, that interpretability is only partial. In the generative variant the next-token prediction is shaped not by the concept bottleneck alone but by the bottleneck together with a parallel layer of unsupervised, uninterpretable neurons added to preserve generation quality. When the dataset labels are reused as concepts the bottleneck can be tiny — 4 neurons for the 4 AG News classes — yet in the authors’ implementation it runs alongside 768 unsupervised ones,<sup>1</sup> so only 4 of the 772 dimensions feeding each token prediction carry interpretable concepts. Those interpretable dimensions remain intervenable nonetheless: because they feed the prediction directly, clamping or amplifying one steers the model’s behaviour predictably — the steerability the bottleneck is built for — even if they carry only a fraction of the signal.

#### 2.2.2.2 Post-hoc Methods

The alternative to *explainable-by-design* models is to leave the trained model untouched and probe it for the concepts it has already come to represent.

**TESTING WITH CONCEPT ACTIVATION VECTORS** T-CAV [19] represents a concept as a direction in a layer’s activation space: given example inputs exhibiting the concept and a set of random ones, it fits a linear classifier to separate them and takes the normal to the decision boundary — the concept activation vector — as the concept’s direction. The model’s sensitivity to the concept is then quantified by the directional derivative of the output along this vector, yielding

<sup>1</sup> The width of the unsupervised layer is fixed at 768 in the authors’ released code (<https://github.com/Trustworthy-ML-Lab/CB-LLMs>); the paper itself leaves it unspecified.

a global statement of how much a given concept influences a class prediction. The concept must still be supplied in advance, as a set of examples.

**SAES AND MECHANISTIC INTERPRETABILITY** A final, unsupervised strand removes even that requirement. Rather than testing for concepts named ahead of time, Sparse AutoEncoders (SAEs) learn a dictionary of concept-like features directly from a model’s internal activations [15], and have become a central tool of Mechanistic Interpretability (MechInt) — the effort to reverse-engineer the computations of neural networks, and of Large Language Models in particular. This unsupervised discovery of interpretable features is treated in full in section 2.3.2.

## 2.3 AUTOENCODERS

An autoencoder is a neural network trained to reconstruct its input through an intermediate latent representation [13, Ch 14]. It consists of two components: an **encoder**  $f_\phi : \mathbb{R}^n \rightarrow \mathbb{R}^m$  that maps an input  $x$  to a latent code  $z = f_\phi(x)$ , and a **decoder**  $g_\theta : \mathbb{R}^m \rightarrow \mathbb{R}^n$  that reconstructs the input as  $\hat{x} = g_\theta(z)$ . Training minimises a reconstruction loss, typically the mean squared error (MSE):

$$\mathcal{L}_{\text{recon}} = \text{MSE}(x, \hat{x}) = \frac{1}{n} \|x - \hat{x}\|_2^2. \quad (2.4)$$

When  $m < n$ , the bottleneck forces the model to learn a compressed representation. When  $m > n$ , the network is overcomplete and additional constraints are required to prevent the trivial solution of learning the identity. In either case, the structure of the latent space is shaped by the choice of architecture and regularisation.

### 2.3.1 VARIATIONAL AUTOENCODERS

Variational AutoEncoders (VAEs) [21, 33] reinterpret the autoencoder as a probabilistic generative model. Rather than mapping an input to a single latent point, the encoder maps it to a *distribution* over latent codes, and the decoder is treated as a generative model that defines a distribution over inputs given a latent code. This probabilistic framing turns the autoencoder into a model from which new data can be sampled, and forces its latent space to have a smooth, continuous structure.

This smoothness is the property that sets a VAE apart from a standard autoencoder. Trained for reconstruction alone, an autoencoder is free to scatter its encodings into well-separated clusters with large gaps between them; those gaps are regions that decode to unrealistic outputs, so the space cannot be sampled or interpolated reliably.




Figure 2.2: Two-dimensional latent spaces learned on MNIST, coloured by digit class, reproduced from Shafkat [37]. Optimising for reconstruction alone (left) lets a standard autoencoder separate the classes into disconnected clusters with gaps that decode to unrealistic samples, whereas the Variational AutoEncoder objective (right) yields a smoother, more continuous arrangement, concentrated around the origin by the prior, that can be sampled and interpolated.

The VAE objective instead pushes encodings to fill the latent space more uniformly and to vary smoothly, so that nearby codes decode to similar inputs and paths between codes stay on the data manifold. Figure 2.2 illustrates the difference on a two-dimensional latent space trained on MNIST.

### 2.3.1.1 Variational Inference

Formally, a VAE models the data through a generative process in which a latent code is first drawn from a prior  $z \sim p(z)$  — typically a standard Gaussian  $p(z) = \mathcal{N}(\mathbf{0}, \mathbf{I})$  — and an observation is then drawn from a decoder distribution  $p_\theta(x | z)$  parameterised by a neural network. Training would ideally maximise the marginal likelihood

$$p_\theta(x) = \int p_\theta(x | z) p(z) dz,$$

but this integral, and the corresponding posterior  $p_\theta(z | x) = p_\theta(x | z)p_\theta(z)/p_\theta(x)$ , are intractable for any non-trivial decoder.

The key idea of Kingma and Welling [21] is to sidestep this intractability with *amortised variational inference*: an encoder network  $q_\phi(z | x)$  is trained to approximate the true posterior, with all data points sharing the same inference parameters  $\phi$ . The encoder outputs the parameters of a diagonal Gaussian,

$$q_\phi(z | x) = \mathcal{N}(z; \mu_\phi(x), \text{diag}(\sigma_\phi^2(x))), \quad (2.5)$$



so that each input is mapped to a mean and a (diagonal) covariance rather than to a single point.

Instead of the intractable likelihood, the [VAE](#) maximises a tractable lower bound on it, the Evidence Lower Bound ([ELBO](#)):

$$\log p_\theta(\mathbf{x}) \geq \underbrace{\mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x} | \mathbf{z})]}_{\text{reconstruction}} - \underbrace{D_{\text{KL}}(q_\phi(\mathbf{z} | \mathbf{x}) \parallel p(\mathbf{z}))}_{\text{regularisation}}. \quad (2.6)$$

The bound follows from the identity

$$\log p_\theta(\mathbf{x}) = \text{ELBO} + D_{\text{KL}}(q_\phi(\mathbf{z} | \mathbf{x}) \parallel p_\theta(\mathbf{z} | \mathbf{x})),$$

together with the non-negativity of the Kullback–Leibler ([KL](#)) divergence; it is tight precisely when the approximate posterior matches the true posterior  $p_\theta(\mathbf{z} | \mathbf{x})$ . Maximising the [ELBO](#) simultaneously pushes the decoder to assign high likelihood to the data (the reconstruction term) and keeps the approximate posterior close to the prior (the [KL](#) divergence term). Negating it yields the training loss:

$$\mathcal{L}_{\text{VAE}} = \mathcal{L}_{\text{recon}} + D_{\text{KL}}(q_\phi(\mathbf{z} | \mathbf{x}) \parallel p(\mathbf{z})). \quad (2.7)$$

When the decoder is Gaussian with fixed isotropic variance, the first term reduces, up to an additive constant, to the [MSE](#) reconstruction loss of eq. (2.4); the second term, being the divergence between two Gaussians, has a closed form. For the standard-normal prior it becomes:

$$D_{\text{KL}}(q_\phi(\mathbf{z} | \mathbf{x}) \parallel p(\mathbf{z})) = \frac{1}{2} \sum_{j=1}^m (\mu_j^2 + \sigma_j^2 - \log \sigma_j^2 - 1),$$

where  $\mu_j$  and  $\sigma_j$  are the components of the encoder’s predicted mean and standard deviation.

Evaluating the reconstruction term requires sampling  $\mathbf{z}$  from  $q_\phi(\mathbf{z} | \mathbf{x})$ , an operation that is not differentiable with respect to the encoder parameters and would therefore block gradient-based training. The [VAE](#) resolves this with the *reparameterisation trick*: the stochasticity is moved to an auxiliary noise variable, and the sample is expressed as a deterministic, differentiable function of the encoder outputs,

$$\mathbf{z} = \boldsymbol{\mu}_\phi(\mathbf{x}) + \boldsymbol{\sigma}_\phi(\mathbf{x}) \odot \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}),$$

where  $\odot$  denotes the element-wise product. Gradients can now flow through  $\boldsymbol{\mu}_\phi$  and  $\boldsymbol{\sigma}_\phi$  while the randomness is confined to  $\boldsymbol{\epsilon}$ , making the whole objective trainable end-to-end by stochastic gradient descent.

### 2.3.1.2 Architectures

There exist several variants that modify the base [VAE](#) to shape the latent space or the form of the latent code; two are particularly relevant here.



$\beta$ -VAE Higgins et al. [14] introduced a single modification: a scalar weight  $\beta$  on the KL term of the objective,

$$\mathcal{L}_{\beta\text{-VAE}} = \mathcal{L}_{\text{recon}} + \beta D_{\text{KL}}(q_{\phi}(z | x) \parallel p(z)). \quad (2.8)$$

Setting  $\beta > 1$  places additional pressure on the approximate posterior to match the factorised prior, which encourages individual latent dimensions to capture statistically independent factors of variation — a property known as *disentanglement*. The standard VAE is recovered at  $\beta = 1$ . Changing  $\beta$  offers a trade-off: a larger value improves the structure and interpretability of the latent space at the cost of reconstruction fidelity.

VQ-VAE van den Oord, Vinyals and Kavukcuoglu [43] replaced the continuous Gaussian latent with a *discrete* one. The encoder output is quantised to the nearest entry of a learned codebook  $\{e_1, \dots, e_K\}$ , so that each latent is represented by a discrete index rather than a continuous vector. This sidesteps the *posterior collapse* that can affect continuous VAEs paired with expressive autoregressive decoders, in which the decoder learns to ignore the latent entirely. Because the nearest-neighbour lookup is non-differentiable, gradients are passed to the encoder with a straight-through estimator and the codebook is learned through dedicated loss terms. Vector-Quantized Variational AutoEncoder (VQ-VAE) underpins many modern generative systems, providing the discrete latent vocabulary over which an autoregressive prior is subsequently trained.

### 2.3.2 SPARSE AUTOENCODERS

Sparse AutoEncoders (SAEs) are one of the main variants of autoencoders. Just like autoencoders, they learn a dictionary of feature directions from the input data, but they typically are overcomplete ( $m \gg n$ ) and their latent representation is forced to be sparse using some kind of sparsity mechanism. Given input  $x \in \mathbb{R}^n$ , the SAE encodes it as:

$$z = \text{ReLU}(W_e(x - b_{\text{pre}}) + b_e), \quad (2.9)$$

where  $W_e \in \mathbb{R}^{m \times n}$  is the encoder weight matrix,  $b_e \in \mathbb{R}^m$  is the encoder bias, and  $b_{\text{pre}} \in \mathbb{R}^n$  is a pre-encoder bias subtracted from the input before encoding. This learned offset allows the encoder to operate on roughly mean-centred activations, improving training stability [4]. The decoder then reconstructs the input as:

$$\hat{x} = W_d z + b_{\text{pre}}, \quad (2.10)$$

where  $W_d \in \mathbb{R}^{n \times m}$  is the decoder weight matrix. Each column  $d_i$  of  $W_d$  is a dictionary atom, representing a candidate feature direction in activation space.

When used in the context of MechInt, the overcompleteness and the sparsity constraint are motivated by the *superposition hypothesis*.




Figure 2.3: Neuron 4e:55, a polysemantic neuron from InceptionV1 which responds to cat faces, fronts of cars, and cat legs. Reproduced from Olah et al. [28].

### 2.3.2.1 The Superposition Hypothesis

A central challenge in [MechInt](#) is that individual neurons in neural networks tend to be *polysemantic*: they activate in response to multiple, seemingly unrelated concepts [10, 28]. An example of this can be seen in fig. 2.3. This runs contrary to the intuition that an interpretable model should have neurons with clean, well-defined roles.

Elhage et al. [10] proposed the *superposition hypothesis* to explain this phenomenon. The core idea is that a network may represent more features than it has dimensions by exploiting the fact that natural inputs are sparse: at any given time, only a small fraction of all possible features are relevant. Under this condition, a network can pack many near-orthogonal feature directions into a lower-dimensional space with limited interference; polysemanticity ceases to be a failure of the model, and instead becomes a rational compression strategy under capacity constraints.

This has a direct implication for interpretability: the natural basis of a model’s activations is not the neuron basis, but a larger, overcomplete set of feature directions embedded within it. Recovering this basis requires moving beyond individual neuron analysis, towards methods that can identify sparse, distributed structure. This is precisely the goal of *sparse dictionary learning* [29], a classical technique from computational neuroscience that [SAEs](#) adapt to the setting of modern language models.

### 2.3.2.2 Architectures

There are many ways to enforce sparsity, each with its pros and cons. Here is a selection of methods that range from the initial [SAE](#) idea to more modern architectures that try to improve on it.

**VANILLA  $\ell_1$**  This is the original architecture proposed by Huben et al. [15]. It uses an  $\ell_1$  penalty<sup>2</sup> on the latent code to encourage most of  $\mathbf{z}$ ’s entries to be exactly zero, so that any given activation is

<sup>2</sup> The  $\ell_p$  norm is defined as  $\ell_p(\mathbf{x}) = \|\mathbf{x}\|_p = (\sum_i |x_i|^p)^{1/p}$ ,  $p \geq 1$ .

reconstructed from a small number of active features. The loss used to train this model is the following:

$$\mathcal{L} = \mathcal{L}_{\text{recon}} + \lambda \|z\|_1. \quad (2.11)$$

The scalar  $\lambda$  controls the sparsity–reconstruction trade-off and must be tuned carefully: too large, and reconstruction degrades; too small, and the learned features lose monosemanticity. A known side effect of  $\ell_1$  regularisation is *shrinkage* [42, 46]: beyond driving inactive features to zero, the penalty also biases active feature magnitudes below their reconstruction-optimal values, causing the model to systematically underestimate the true activation of each feature [31, 46]. Constraining the decoder columns to unit  $\ell_2$  norm after each gradient step partially mitigates this, by preventing the model from reducing the  $\ell_1$  penalty through the trivial reparameterisation of scaling down encoder outputs and compensating with larger decoder norms [31]. There is, however, a residual gradient-level bias that is structurally entangled with the  $\ell_1$  penalty and cannot be resolved with this trick.

**TOPK SAES** Gao et al. [12] proposed replacing the  $\ell_1$  penalty with a hard sparsity constraint, implemented via a TopK activation function that retains only the  $k$  largest pre-activations and zeros the rest. This directly controls the number of active features per forward pass, eliminating the need to tune  $\lambda$  and removing the shrinkage bias entirely. The full training objective is:

$$\mathcal{L} = \mathcal{L}_{\text{recon}} + \alpha \mathcal{L}_{\text{aux}}, \quad (2.12)$$

where the auxiliary loss  $\mathcal{L}_{\text{aux}}$  addresses the problem of dead latents, i. e., features that cease to activate entirely during training. Latents are flagged as dead if they have not activated for a predetermined number of tokens. Given the reconstruction residual  $e = x - \hat{x}$ , the top- $k_{\text{aux}}$  dead latents are selected and used to reconstruct  $e$ ; the resulting reconstruction error is  $\mathcal{L}_{\text{aux}}$ . This gives dead latents a gradient signal proportional to how much of the residual they could explain, encouraging them to activate rather than remain permanently dormant. Gao et al. [12] demonstrate that TopK SAEs achieve better reconstruction at equivalent sparsity levels compared to  $\ell_1$  baselines, making them a strong and widely used alternative.

**GATED SAES** Rajamanoharan et al. [31] introduced an architecture that separates the decision of whether a feature is active from the decision of how strongly it activates. A dedicated gating pathway produces a binary mask via a Heaviside step function (approximated during training with a straight-through estimator) which is applied element-wise to the encoder output. This architectural separation allows the model to learn feature presence and magnitude more independently, addressing shrinkage from a structural rather than a loss-based perspective.

**JUMPRELU SAES** Rajamanoharan et al. [32] proposed replacing the standard ReLU with a JumpReLU activation, which is zero below a learned per-feature threshold  $\theta_i$  and equal to the pre-activation above it:

$$\text{JumpReLU}_{\theta_i}(z_i) = z_i \cdot \mathbb{1}[z_i > \theta_i]. \quad (2.13)$$

The key advantage of this parameterisation is that it provides a dedicated, per-feature learnable threshold, which allows the  $\ell_0$  norm<sup>3</sup> of the feature activations —  $\sum_i H(\pi_i - \theta_i)$ , where  $\pi_i$  are the encoder pre-activations and  $H(x)$  is the Heaviside step function — to be trained directly via a straight-through estimator applied to the threshold parameter  $\theta_i$ . Because each threshold can be positioned near the actual boundary between active and inactive pre-activations for its feature, the gradient signal is stable and controllable, making  $\ell_0$  training practical in a way that would not be possible with a fixed threshold. Training with  $\ell_0$  rather than  $\ell_1$  eliminates shrinkage entirely, since  $\ell_0$  penalises only the number of active features and not their intensity. JumpReLU SAEs achieve comparable performance to TopK SAEs.

### 2.3.2.3 Applications to MechInt

SAEs find their primary application in interpreting the internal activations of neural networks, and MechInt remains their most prominent application, making it the focus of this section. SAEs can be trained on activations from any layer or component of a transformer (section 2.1.2); common targets include the residual stream at various depths, MLP layer outputs, and attention head outputs [4, 15]. The resulting dictionary then supports the applications described below.

**FEATURE DISCOVERY** The primary use of SAEs is *feature discovery*: training an SAE on a model’s internal activations and examining the resulting dictionary atoms to identify human-interpretable concepts. The underlying hypothesis is that each atom, active for only a semantically coherent set of inputs, corresponds to a monosemantic feature of the model’s computation — a clean unit of meaning that the model had obscured in the neuron basis. In practice these features often correspond to recognisable concepts: Bricken et al. [4] trained an SAE on the MLP of a one-layer transformer and found features responsive to specific tokens, syntactic roles, and semantic categories, while at production scale Templeton et al. [41] report features for famous people, countries, emotions, or even biases (like racism, sexism, hatred) and sarcastic praise. These results provide strong empirical support for the superposition hypothesis and suggest that SAEs can serve as a practical lens through which to study what information LLMs represent.

<sup>3</sup> The  $\ell_0$  norm is a pseudo-norm defined as the number of non-zero vector components:  $\ell_0(\mathbf{x}) = \sum_i \mathbb{1}[x_i > 0]$ .

**STEERING** The features recovered by an SAE are not merely descriptive: because each corresponds to a direction in activation space, its activation can be clamped to an arbitrary value during a forward pass to causally *steer* the model’s behaviour. Templeton et al. [41] demonstrated this by amplifying a single feature, causing the model to compulsively reference the Golden Gate Bridge regardless of the input. Such interventions show that the discovered features are causally implicated in the model’s computation rather than merely correlated with its inputs, and open a path to controllable generation and targeted safety interventions.

**CIRCUIT ANALYSIS** Beyond cataloguing individual features, SAEs can serve as the building blocks for reconstructing the *circuits* a model uses to perform a computation, expressing the interactions between features across layers as an interpretable causal graph [25]. More recent work favours *transcoders* — a variant trained to approximate the input–output behaviour of an MLP sublayer with a sparse hidden representation, rather than to reconstruct a single activation — which should disentangle the model’s computation more cleanly and yield sparser, more faithful circuits [1, 8].

#### 2.3.2.4 Evaluation

Assessing the quality of learned features is non-trivial, as interpretability is inherently subjective. Several proxy metrics are used in practice. *Reconstruction quality* is judged by its effect on the model rather than by the raw reconstruction error: the SAE’s output is substituted for the original activations, the model is run forward, and the resulting increase in its next-token CE loss is measured [12]. A faithful SAE leaves this loss almost unchanged, indicating that it has retained the information the model actually uses. *Sparsity* is quantified via the average  $\ell_0$  norm of the latent codes, i.e., the mean number of active features per token. *Automated interpretability* pipelines use a second language model to generate and score natural-language descriptions of individual features [2], providing a scalable if imperfect proxy for human judgement.

## CONCLUSIONS

This chapter has assembled the three strands the thesis builds on: the transformer and its residual stream (section 2.1), whose dense, polysemantic activations are the object to be decomposed; the shift in explainability from attribution to concepts (section 2.2), and the concept-based methods that describe a model’s computation in human-meaningful terms; and the autoencoder family (section 2.3), from

the variational and discrete latents of [VAEs](#) and [VQ-VAEs](#) to the overcomplete, sparse dictionaries of [SAEs](#).

A single limitation runs through these concept-based methods. Whether a concept is a unit in the bottleneck of [a CBM](#) or an atom in the dictionary of [an SAE](#), it is encoded as a single scalar and can vary only in magnitude. This is what forces the sparsity–reconstruction trade-off at the heart of the [SAE](#) literature — a faithful reconstruction needs many active concepts, an interpretable decomposition only few — and the architectures of [section 2.3.2.2](#) all work within it. [CEMs](#) are the exception already met: by giving each concept a vector embedding rather than a scalar, they recover the capacity a scalar bottleneck gives up. [Section 3.3](#) carries that idea into the unsupervised setting of activation decomposition — combining the vector-valued concepts of [CEMs](#) with the variational inference and discrete latents developed in [section 2.3](#) — and derives the resulting model from first principles.

## METHODS

---

### 3.1 AIM OF THE WORK

SAEs have become the standard tool for decomposing the activations of an LLM into a dictionary of concept-like features (section 2.3.2). Their reach is bounded, though, by a structural trade-off between sparsity and reconstruction [12]: because each concept enters the reconstruction through a single scalar coefficient, reconstructing an activation faithfully tends to require many concepts active at once, whereas the interpretability that motivates the decomposition rests on keeping only a few active per token [4, 15]. Pushed to a useful sparsity, the reconstruction degrades far enough that the reconstructed activation leaves the data manifold. This is more than an aesthetic flaw: any downstream computation is then performed on an out-of-distribution signal even before any intervention, and when a feature is amplified or ablated the resulting change in behaviour conflates the intended edit with this uncontrolled reconstruction error, so the very analyses the decomposition is built for become less reliable than they appear.

The limitation lies in the representation rather than the training: a scalar can record only how strongly a concept is present. This thesis asks whether giving each concept a *vector* embedding relaxes the trade-off. The idea is familiar from CEMs (section 2.2.2), which replace the scalar concept bottleneck of a CBM with per-concept embeddings and so recover the accuracy a narrow bottleneck gives up [11]. CEMs, however are supervised, trained against concept labels. The aim here is to carry the concepts-as-vectors idea into the *unsupervised* setting of activation decomposition, and to do so from first principles, defining a probabilistic generative model whose inference procedure and training objective follow from the model itself rather than being assembled by hand. The starting point is the Learnable Concept-based Model (LCBM) of De Santis et al. [6], an unsupervised concept-based image classifier whose encoder produces a vector embedding per concept, gates each concept by comparing that embedding to a learnable prototype, and reconstructs the input and simultaneously classifies it from the embeddings of the active concepts. LCBM supplied the idea rather than the architecture: this thesis keeps its core structure — concepts as embeddings gated against prototypes and decoded by reconstruction — but rebuilds it from the ground up as a fully variational generative model, giving the concept embeddings a probabilistic latent that



LCBM leaves deterministic and dropping the supervised classification task, so that the model reconstructs activations from no labels at all.

The Variational Auto Embedding Encoder (VAEE) sits at the intersection of three of the families surveyed in chapter 2. From CEMs it takes the representation of a concept as a high-dimensional embedding rather than a scalar; from VAEs it takes a probabilistic latent and amortised variational inference (section 2.3.1); and from VQ-VAEs it takes the idea of a discrete latent over a learned dictionary of vectors (section 2.3.1.2). The two differ in what is made discrete: a VQ-VAE quantises each input to a single codebook entry, so its latent is a categorical index and the vector passed to the decoder is one of finitely many. The VAEE instead keeps its embeddings continuous, its only discrete latent being the binary mask  $c$ , which decides *which* prototype-anchored concepts are active rather than what their embeddings are. It is in this sense a probabilistic, multilabel counterpart to vector quantisation: independent Bernoulli gates over continuous embeddings in place of a single hard codebook lookup, keeping the expressivity of a continuous latent while holding the active set sparse.

### 3.2 NOTATION

The VAEE operates on a single input vector  $x \in \mathbb{R}^n$  — in the experiments of chapter 4, a residual stream state of an LLM (section 2.1.2) — and decomposes it into a dictionary of  $K$  concepts. Where an SAE represents each concept by a single scalar coefficient (section 2.3.2), the VAEE represents it by a  $d$ -dimensional embedding. The symbols used throughout the chapter are:

- $x \in \mathbb{R}^n$ : the input vector, to be reconstructed;
- $K$ : the number of concepts in the dictionary;
- $d$ : the dimension of each concept embedding;
- $c = (c_1, \dots, c_K) \in \{0, 1\}^K$ : the binary *concept-activation mask*, where  $c_k = 1$  marks concept  $k$  active for the given input;
- $z_k \in \mathbb{R}^d$ : the *embedding* of concept  $k$  for the given input, collected into  $z = [z_1, \dots, z_K] \in \mathbb{R}^{Kd}$ ;
- $h_k \in \mathbb{R}^d$ : the learned *prototype* of concept  $k$ , a global parameter shared across all inputs;
- $\mu_k(x) \in \mathbb{R}^d$ : the encoder’s output for concept  $k$ , the mean of the embedding  $z_k$ ;
- $\alpha_k(x) \in [0, 1]$ : the probability that concept  $k$  is active for input  $x$ .

As in section 2.3,  $f_\phi$  denotes the encoder and  $g_\theta$  the decoder, with parameters  $\phi$  and  $\theta$  respectively.



The word *prototype* needs a caveat, since section 2.2.1.3 used it in a slightly different sense. There a prototype was an *input instance* — an actual data point, typical of a region of the data, chosen to summarise it by inspection. The VAAE’s prototypes are similar in their function, as they act as a centroid of sorts for a group of input data, but instead of living in the input space they live in latent space.

### 3.3 MODEL DEVELOPMENT

#### 3.3.1 PROBABILISTIC FORMULATION

The VAAE is defined as a probabilistic generative model of an observation  $x$ , built so that a sparse set of concept embeddings accounts for it.

##### 3.3.1.1 Generative model

The VAAE places a joint distribution over the mask  $c$ , the embeddings  $z$ , and the observation  $x$  that factorises along the generative flow  $c \rightarrow z \rightarrow x$ ,

$$p(c, z, x) = p(c) p(z | c) p_\theta(x | z). \quad (3.1)$$

Its three factors are specified in turn. First, a binary mask decides whether the concept is present: each  $c_k$  is drawn independently from a Bernoulli with a fixed prior activation rate  $\pi$ , shared across concepts,

$$p(c_k) = \text{Bern}(c_k; \pi). \quad (3.2)$$

The mask then represents a *Gated Gaussian* prior over the embedding  $z$ : an inactive concept ( $c_k = 0$ ) draws its embedding from a Gaussian at the origin, an active one ( $c_k = 1$ ) from a Gaussian centred on the concept’s prototype  $h_k$ ,

$$p(z_k | c_k) = \mathcal{N}(z_k; c_k h_k, \sigma_0^2 \mathbb{I}), \quad (3.3)$$

both with a fixed isotropic variance  $\sigma_0^2$ . The prototype  $h_k$  thus has a precise meaning: it is the mean embedding of concept  $k$  when the concept is active — its canonical location in embedding space. Finally, the decoder  $g_\theta$  maps the embeddings to the observation through a Gaussian likelihood centred on its output,

$$p_\theta(x | z) = \mathcal{N}(x; g_\theta(z), \mathbb{I}). \quad (3.4)$$

Figure 3.1 shows this structure as a probabilistic graphical model: the mask fixes where each embedding sits, and the observation follows from the embeddings.




Figure 3.1: The VAAE generative model as a probabilistic graphical model. For each of the  $K$  concepts, a binary mask  $c_{i,k}$  (drawn from a Bernoulli with rate  $\pi$ ) gates a Gaussian embedding  $z_{i,k}$  centred either at the origin or at the learned prototype  $h_k$ ; the gated embeddings generate the observation  $x_i$ . Shaded nodes are observed and unshaded latent, bare symbols are global parameters; the solid plate replicates over concepts and the dashed plate over the  $N$  observations.

### 3.3.1.2 Inference

Exact inference is intractable: it would require marginalizing over the continuous embeddings  $z$  and summing over all  $2^K$  possible discrete masks configurations for  $c$ :

$$p(x) = \sum_c \int p(x, c, z) dz.$$

The VAAE instead uses amortised variational inference (section 2.3.1), approximating the true posterior with a *mean-field* family that factorises over both latents and across concepts,

$$q_\phi(c, z | x) = \prod_{k=1}^K q_\phi(z_k | x) q_\phi(c_k | x), \quad (3.5)$$

with an encoder  $f_\phi$  that predicts the parameters of each factor, shared across all data points. The encoder predicts a mean embedding  $\mu_k(x)$  for each concept; the embedding posterior is Gaussian about this mean, with the posterior variance frozen to the prior's  $\sigma_0^2$ ,

$$q_\phi(z_k | x) = \mathcal{N}(z_k; \mu_k(x), \sigma_0^2 \mathbb{I}), \quad (3.6)$$

or, using the reparameterisation trick,

$$z_k = \mu_k(x) + \sigma_0 \epsilon, \quad \epsilon \sim \mathcal{N}(\mathbf{0}, \mathbb{I}).$$

The activation probability is not predicted by a separate head but *derived* from the predicted embedding location  $\mu_k(x)$ : a concept should

be judged active precisely when its predicted embedding lands near its prototype. Because the gate reads off the mean  $\mu_k$  rather than a sampled embedding, the mask depends on  $x$  alone, as the mean-field posterior requires. As shown in appendix A.1, this gate is the optimal mean-field mask factor under the Gated Gaussian prior: a sigmoid of the similarity between the predicted embedding and the prototype,

$$\alpha_k(x) = \text{sigmoid} \left( \frac{\text{sim}(\mu_k(x), h_k)}{\tau} \right), \quad (3.7)$$

scaled by an inverse temperature  $1/\tau$  — a single learnable scalar, initialised to  $\tau = 0.1$ , in the manner of the CLIP logit scale. This activation probability parameterises the Bernoulli posterior over the mask,

$$q_\phi(c_k | x) = \text{Bern}(c_k; \alpha_k(x)). \quad (3.8)$$

Sampling the mask directly from this Bernoulli is not differentiable with respect to  $\alpha_k$ , which would block the gradient from reaching the encoder through the gate. During training the draw is therefore replaced by the reparameterisable *Gumbel-Sigmoid* relaxation, the binary, multilabel analogue of the Gumbel-Softmax (Concrete) distribution [16, 24]: rather than relaxing a single mutually-exclusive choice over categories, it applies an independent relaxed Bernoulli to each concept, so any subset of concepts may be active at once. A Bernoulli draw can be written as  $\mathbb{1}[u < \alpha_k(x)] = \mathbb{1}[\text{logit} \alpha_k(x) > g]$  with logistic noise  $g = \text{logit}(u)$ ,  $u \sim \text{Uniform}(0, 1)$ ; <sup>1,2</sup> replacing the indicator with a temperature-controlled sigmoid gives the differentiable surrogate

$$\tilde{c}_k = \text{sigmoid} \left( \frac{\text{logit} \alpha_k(x) - g}{\tau_g} \right). \quad (3.9)$$

The relaxation temperature  $\tau_g$  sets the sharpness: as  $\tau_g \rightarrow 0$  the sigmoid tends to the indicator and  $\tilde{c}_k$  recovers an exact  $\{0, 1\}$  Bernoulli sample, while a larger  $\tau_g$  gives smoother values that keep the gradient well-behaved. At inference the stochastic draw is dropped and both latents are replaced by their posterior means — the embedding  $z_k$  by  $\mu_k$  and the mask  $c_k$  by its expectation  $\alpha_k(x)$ . The gate is then nominally soft, but in practice effectively hard: the binarisation pressure of the training objective (section 3.3.2) drives the learned activations to a strongly bimodal distribution concentrated near 0 and 1.

One subtlety remains. The embedding posterior of eq. (3.6) does not depend on the mask, since the encoder produces  $\mu_k$  from  $x$  alone, so a sampled  $z_k$  would pass signal to the decoder regardless of whether the concept is active. To enforce that inactive concepts contribute exactly nothing, a multiplicative gate is applied before decoding,

$$\hat{x} = g_\theta(c \odot z), \quad (c \odot z)_k = c_k z_k. \quad (3.10)$$

<sup>1</sup> The logit transform is defined as  $\text{logit}(s) = \log \frac{s}{1-s}$ .

<sup>2</sup> This logistic noise is the difference of two standard Gumbel variates, which is what ties the relaxation to the Gumbel-Softmax family.

The gate is empirically load-bearing: without it, the soft penalty on inactive embeddings derived in section 3.3.2 is too weak to drive them to zero, because the reconstruction gradient keeps inflating them. It also exposes the mask  $c$  as a clean handle for concept-level interventions (section 3.3.3).

### 3.3.1.3 Similarity

The gate of eq. (3.7) is written for a generic similarity  $\text{sim}(\cdot, \cdot)$ , for which three choices are considered:

$$\text{sim}(\mu_k, h_k) = \begin{cases} \frac{\mu_k \cdot h_k}{\|\mu_k\| \|h_k\|} & \text{(cosine similarity)} \\ \mu_k \cdot h_k & \text{(inner product)} \\ b - \|\mu_k - h_k\| & \text{(negative Euclidean)} \end{cases} \quad (3.11)$$

Because the embeddings are decoded unnormalised, the choice determines whether the gate is coupled to the embedding’s magnitude. *Cosine similarity* depends only on the angle between  $\mu_k$  and  $h_k$ : the angle controls the discrete gate while the magnitude is left free to carry information to the decoder. The plain inner product is unbounded and entangles gating with magnitude (the network can inflate  $\|\mu_k\|$  to force a concept on) while the negative Euclidean distance fits the Gaussian prior most directly but requires learning the radius offset  $b$ . Cosine similarity is used throughout unless stated otherwise.

### 3.3.2 TRAINING OBJECTIVE

The VAAE is trained by maximising a lower bound on the marginal log-likelihood of the data, just as the VAE of section 2.3.1. Under the mean-field approximation of eq. (3.5) the ELBO decomposes into three analytic components:

$$\begin{aligned} \log p_\theta(\mathbf{x}) \geq & \underbrace{\mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x} | \mathbf{z})]}_{\text{reconstruction}} + \\ & - \sum_{k=1}^K \underbrace{D_{\text{KL}}(q_\phi(c_k | \mathbf{x}) \parallel p(c_k))}_{\text{mask-prior KL}} + \\ & - \sum_{k=1}^K \underbrace{\mathbb{E}_{q_\phi(c_k|\mathbf{x})} [D_{\text{KL}}(q_\phi(\mathbf{z}_k | \mathbf{x}) \parallel p(\mathbf{z}_k | c_k))]}_{\text{conditional-embedding KL}}. \end{aligned} \quad (3.12)$$

The full ELBO derivation is given in appendix A.2. The three terms read directly off the generative model. Under the Gaussian decoder of eq. (3.4) the reconstruction term reduces, up to an additive constant, to the MSE of eq. (2.4); the mask-prior KL pulls each concept’s activation toward the Bernoulli prior rate  $\pi$ ; and the conditional-embedding KL

keeps each embedding near its prototype when the concept is active and near the origin when it is not.

### 3.3.2.1 From bound to objective

Turning the [ELBO](#) into a practical loss takes a few deliberate departures from a strict negated bound, each motivated by interpretability or training stability rather than mathematical necessity.

**BATCH-WISE SPARSITY** Conventional sparsity penalties — the  $\ell_1$  and TopK constraints of section [2.3.2](#) — act along the concept axis, bounding how many concepts a *single sample* may activate. Applied per sample, the mask-prior [KL](#) would do the same, pulling every input’s  $\alpha_k(x)$  toward  $\pi$  independently and so holding every sample to one fixed activation budget. We instead constrain sparsity along the *batch* axis [\[6\]](#): rather than how many concepts a sample uses, we limit how often a *given* concept fires across a batch. The per-sample term is replaced by a single [KL](#) between Bernoullis that matches each concept’s batch-average activation  $\bar{\alpha}_k = \frac{1}{B} \sum_i \alpha_k(x_i)$  to the prior rate,

$$\beta \sum_{k=1}^K D_{\text{KL}}(\bar{\alpha}_k \parallel \pi),$$

with  $B$  the batch size and  $\beta$  the term’s weight; matching  $\bar{\alpha}_k$  to  $\pi$  lets concept  $k$  fire on about a  $\pi$  fraction of the batch.

A per-sample budget is not as good a target because inputs differ in complexity: some are explained by a few generative features, others by many. After shuffling, a batch mixes both, and it is unlikely that many of its samples all call for the same concept. Constraining each concept’s firing rate rather than each sample’s count therefore leaves the per-sample count free, letting the model spend many concepts on a complex input and few on a simple one, while still keeping any single concept from firing everywhere or dying out across the dataset without the need of an auxiliary loss to revive ded latents.

**ENTROPY BINARISATION** Batch-wise sparsity does not by itself force per-sample decisions to be sharp: activations clustered around  $\pi$  would satisfy it while remaining soft. We add a per-sample binary-entropy penalty,

$$\lambda_{\text{ent}} \mathcal{H}(\alpha_{i,k}) = -\lambda_{\text{ent}} [\alpha_{i,k} \log \alpha_{i,k} + (1 - \alpha_{i,k}) \log(1 - \alpha_{i,k})],$$

which is maximal at  $\alpha = \frac{1}{2}$  and zero at the binary extremes, so minimising it drives each gate toward  $\{0, 1\}$  and aligns the soft training-time activation with the binary mask of the generative model. This is a practical add-on, not a rearranged [ELBO](#) term.

**STOP-GRADIENT ON THE ACTIVATION** The conditional-embedding [KL](#) couples the activation to the embeddings: in closed form (appendix [A.2](#)) it reads  $\alpha_{i,k} \|\boldsymbol{\mu}_{i,k} - \mathbf{h}_k\|_2^2 + (1 - \alpha_{i,k}) \|\boldsymbol{\mu}_{i,k}\|_2^2$ , which the network could shrink by driving  $\alpha_{i,k} \rightarrow 0$  irrespective of reconstruction. A stop-gradient  $\text{sg}(\cdot)$  operator on  $\alpha$  inside this term removes that shortcut, so the activation is shaped only by the reconstruction and sparsity terms.

#### DECODER ORTHOGONALITY

**SCALE-INVARIANT REDUCTIONS** Each term is averaged over its dimensions — the batch  $B$ , the concepts  $K$ , and the relevant vector dimension (the input  $n$  or the embedding  $d$ ) — so a single set of weights  $\gamma, \beta, \lambda_{\text{ent}}$  transfers across architectures of different sizes without rescaling.

##### 3.3.2.2 Practical loss

Collecting these choices and substituting the closed-form conditional-embedding [KL](#) of appendix [A.2](#) gives the objective minimised in training,

$$\begin{aligned} \mathcal{L} = & \frac{1}{Bn} \sum_{i=1}^B \|\mathbf{x}_i - \hat{\mathbf{x}}_i\|_2^2 + \frac{\beta}{K} \sum_{k=1}^K D_{\text{KL}}(\tilde{\alpha}_k \parallel \pi) + \frac{\lambda_{\text{ent}}}{BK} \sum_{i=1}^B \sum_{k=1}^K \mathcal{H}(\alpha_{i,k}) \\ & + \frac{\gamma}{BKd} \sum_{i=1}^B \sum_{k=1}^K \left( \text{sg}(\alpha_{i,k}) \|\boldsymbol{\mu}_{i,k} - \mathbf{h}_k\|_2^2 + (1 - \text{sg}(\alpha_{i,k})) \|\boldsymbol{\mu}_{i,k}\|_2^2 \right), \end{aligned} \quad (3.13)$$

where  $\hat{\mathbf{x}}_i = g_\theta(\mathbf{c}_i \odot \mathbf{z}_i)$  is the gated reconstruction of eq. [\(3.10\)](#) and  $\boldsymbol{\mu}_{i,k} = \boldsymbol{\mu}_k(\mathbf{x}_i)$ . The four terms are, in order, the reconstruction error, the batch-wise sparsity [KL](#), the entropy binarisation penalty, and the conditional-embedding [KL](#) (active concepts pulled to their prototype, inactive ones to the origin).

##### 3.3.3 ARCHITECTURES

## CONCLUSIONS

## EXPERIMENTS

---

Chapter 3 derived the VAAE as a probabilistic generative model but left its empirical behaviour open. This chapter puts the architecture’s central claim to the test — that representing a concept by a vector rather than a scalar relaxes the sparsity–reconstruction trade-off that bounds SAEs (section 3.1) — on one application of it: the decomposition of LLM activations.

Section 4.1 describes the activation dataset, drawn from the residual stream of a pretrained LLM. Section 4.2 introduces the SAE baseline against which the VAAE is compared at matched parameter count. Section 4.3 sets out the evaluation metrics — a measure of reconstruction fidelity and two of sparsity. Section 4.4 reports the training setup, and section 4.5 presents the quantitative, qualitative, and ablation results.

*Review this paragraph after whole chapter is written.*

### 4.1 EXPERIMENT SETUP

The VAAE operates on a single activation vector  $x \in \mathbb{R}^n$  (section 3.2), so evaluating it on a language model requires a corpus of such vectors. These were extracted from the residual stream of Mistral 7B v0.1 [17], a pretrained decoder-only LLM, run over SetFit’s version<sup>1</sup> of the Stanford Sentiment Treebank (SST-2) sentiment classification dataset [39].

SST-2 is a dataset of movie reviews, each with a label (positive/negative) attached. For the purposes of the experiments we were only interested in the reviews as a corpus of natural language, so the labels were discarded. SST-2 is a small corpus: large enough to exercise reconstruction, but, in hindsight, probably too narrow for a large and varied concept dictionary to emerge, a limitation returned to when the results are interpreted (section 4.5). This problem was exacerbated by the choice of SetFit’s version instead of the original version of the dataset, since the former contains only 6 920 reviews in the training set, while the latter has 67 300 reviews in the training set. The total number of tokens in each set is shown in table 4.1.

Each review was passed through the model and its residual-stream state was read off after the final transformer block, just before the classifier head. This yielded the residual-stream representation of every token in the sentence. The resulting activations have dimension  $n = 4\,096$ , i. e., Mistral’s residual-stream dimension.

---

<sup>1</sup> Available for download at <https://huggingface.co/datasets/SetFit/sst2>.

| SET        | NO. REVIEWS | NO. TOKENS |
|------------|-------------|------------|
| Training   | 6 920       | 179 146    |
| Validation | 872         | 22 887     |

Table 4.1: The distribution of reviews and tokens across the training and validation sets of SetFit/[SST-2](#).

## 4.2 BASELINES

The [VAEE](#) is compared against a TopK [SAE](#) (section 2.3.2) [12], the state-of-the-art model for [MechInt](#). Its top- $k$  activation fixes the number of active features per sample directly through  $k$ , so the reconstruction error can be read against a sparsity level that is set rather than induced by a penalty weight. Sweeping  $k$  traces the baseline’s sparsity–reconstruction frontier, against which the [VAEE](#)’s operating point is placed.

The two models are matched on the size of their encoder and decoder rather than on dictionary size, so that any difference reflects how each represents a concept and not how much capacity it is given. Both map the activation through an encoder into a latent code and back through a decoder; that code has  $m$  entries for an [SAE](#) and  $K \cdot d$  for a [VAEE](#) — its  $K$  concepts of embedding dimension  $d$ , written  $K \times d$  — and the encoder and decoder weight matrices are sized to match. Matching therefore sets  $K \cdot d = m$ : the  $128 \times 64$  [VAEE](#) is paired with the  $m = 8\,192$  [SAE](#),  $64 \times 64$  with  $m = 4\,096$ , and  $32 \times 64$  with  $m = 2\,048$ .<sup>2</sup> The [SAE](#)’s dictionary is correspondingly far larger than the [VAEE](#)’s concept count — each [VAEE](#) concept spends  $d$  parameters where an [SAE](#) atom spends one.

## 4.3 EVALUATION METRICS

Two properties are measured: how faithfully a model reconstructs an activation, and how sparse the decomposition it uses to do so is. Reconstruction is reported as a single error metric; sparsity is reported in two ways, because the [VAEE](#) and the [SAE](#) do not count activity in the same unit.

<sup>2</sup> This equalises the encoder and decoder weights, which dominate the parameter count ( $\approx 2n \cdot Kd$  for both models). Beyond them the [VAEE](#) carries a handful of parameters with no [SAE](#) counterpart — its  $K$  learnable prototypes, of dimension  $d$ , and the gate temperature — so exact matching of the total would pair the three [VAEE](#) sizes with [SAEs](#) of  $m = 8\,194$ ,  $4\,097$ , and  $2\,049$  rather than the powers of two used here. The surplus is one or two dictionary atoms, under 0.05% of the count, so the totals agree to well within a percent.



## 4.3.1 RECONSTRUCTION

Reconstruction fidelity is measured by the [MSE](#) between an activation and its reconstruction, averaged across the whole dataset. This is the same quantity defined in eq. (2.4). It is the natural choice here: under the Gaussian decoder of eq. (3.4) the [VAEE](#)’s reconstruction term reduces to exactly this error (section 3.3.2), and it is the standard reconstruction measure for [SAEs](#), so the two models are scored on the objective each is trained to optimise.<sup>3</sup>

Evaluating reconstruction is important, as a model that has low reconstruction error is a model that has learned *representative* features of the input data, and concepts are representative features by definition.

## 4.3.2 SPARSITY

A low active-concept count is not the same as a sparse code. The [VAEE](#) targets the number of concepts active per sample, not the fraction of its dictionary left idle, so at higher  $\kappa$  it may fire a sizable share of its concepts without contradicting its aim. Dictionary-fraction sparsity is the [SAE](#)’s design principle, not the [VAEE](#)’s; the quantity that bears on the comparison is the absolute number of active units, reported below in two ways.

**NUMBER OF ACTIVE CONCEPTS  $\kappa$**  From an interpretability standpoint, the central sparsity metric is the number of active concepts: this is the number of units a human must understand in order to explain the input instance. For a TopK [SAE](#) this corresponds to its  $k$  hyperparameter by construction;<sup>4</sup> for a [VAEE](#),  $\kappa = \#\{k : \alpha_k(x) \geq 1/2\}$ , the count of concepts whose posterior gate exceeds a threshold. Although the inference gate is continuous (section 3.3.1.2), the threshold introduces no real arbitrariness, because the learned activations are strongly bimodal — concentrated near 0 and 1, with little mass in between, as quantified in section 4.5 — so the count is insensitive to the exact cut-off, and 1/2 was chosen as a sensible threshold for an activation probability.

**$\ell_0$  COUNT** The  $\ell_0$  pseudo-norm represents the amount of information the decoder has access to in order to reconstruct the input. This amount corresponds to the number of non-zero scalars it can reconstruct from. For a TopK [SAE](#), this number is its  $k$  hyperparameter

<sup>3</sup> The two reconstruction checks that are standard for [SAEs](#) spliced back into a language model — the [CE](#) loss and perplexity of the patched model — are not reported for this comparison, as they were not run for the [VAEE](#); [MSE](#) is the only reconstruction metric on which both models are evaluated here.

<sup>4</sup> Technically, TopK selects the  $k$  largest activations, but it is possible that after the ReLU only  $k' < k$  activations are non-zero, so  $\ell_0 = \kappa = k' < k$ . In practice, this almost never happens, but the  $\ell_0$  is recomputed every time for maximum accuracy nonetheless.

| METRIC                     | SAE                             | VAEE                                      |
|----------------------------|---------------------------------|-------------------------------------------|
| $\kappa$ (active concepts) | $\#\{k : z_k(\mathbf{x}) > 0\}$ | $\#\{k : \alpha_k(\mathbf{x}) \geq 1/2\}$ |
| $\ell_0$ (active neurons)  | $\kappa$                        | $\kappa \cdot d$                          |

Table 4.2: Sparsity metric formulations for SAE and VAEE.

by construction;<sup>5</sup> for a VAEE it is the number of active concepts  $\kappa$  multiplied by the embedding dimension of each concept  $d$ .

For a fair comparison it is important to consider both metrics, since a VAEE that activates fewer *concepts* than an SAE may still activate more *neurons* per sample.  $\kappa$  is therefore the fair unit for comparing the two models on interpretability, while  $\ell_0$  is the unit of choice for comparing decoder capacity. Table 4.2 recaps the formulas used to compute  $\kappa$  and  $\ell_0$  for both models.

#### 4.4 MODEL TRAINING

Unless stated otherwise, the VAEE uses the shallow configuration of section 4.5.3: a single linear-plus-GELU layer for the encoder  $f_\phi$  and a linear decoder  $g_\theta$ , the stacked topology, and cosine similarity for the gate (eq. (3.11)). The embedding posterior variance is fixed at  $\sigma_0 = 0.1$  during training; at inference time the sampling is dropped and the embedding  $z_k$  is set to its mean value  $\mu_k$ . The Gumbel-Sigmoid relaxation (eq. (3.9)) uses temperature  $\tau_g = 0.5$  during training; at inference the relaxation is dropped and the gate  $c_k$  set to its posterior mean  $\alpha_k$ .

The VAEE is trained to minimise the objective of eq. (3.13) — reconstruction error, batch-wise sparsity KL, entropy binarisation, and the conditional-embedding KL. The loss weights are set to  $\beta = 4$ ,  $\lambda_{\text{ent}} = \gamma = 1$ . The Bernoulli prior rate  $\pi$  is the VAEE’s sparsity dial, swept to trace its frontier, but it controls sparsity only indirectly, and  $\tilde{\kappa} = \pi K$  is a target the optimisation is biased toward, not a hard constraint. The final per-sample count  $\kappa$  runs higher than the nominal  $\tilde{\kappa}$ , since reconstruction drives  $\kappa$  towards higher values.

The TopK SAE baseline is trained with its standard objective — reconstruction error under the top- $k$  activation plus auxiliary loss — on the same activations.

Both models were trained using the Adam optimiser [20] with a learning rate of  $10^{-4}$  for 3 000 epochs and an early stopping patience of 500 epochs.

<sup>5</sup> See footnote 4.




Figure 4.1: Sparsity–reconstruction frontier in active concepts  $\kappa$ , the interpretability-relevant unit; each model sweeps its sparsity dial at matched parameter count. The **VAEE** frontiers lie below every **SAE** frontier — it matches the best **SAE**’s reconstruction at far fewer active concepts.

## 4.5 RESULTS

### 4.5.1 QUANTITATIVE RESULTS

Figure 4.1 plots validation **MSE** against the number of active concepts  $\kappa$ , sweeping each model’s sparsity dial. The models are matched on parameter count (section 4.2). Across the swept range every **VAEE** frontier lies below every **SAE** frontier: at any given concept count it reconstructs more faithfully, it matches the best **SAE**’s reconstruction using roughly an order of magnitude fewer active concepts, and at the dense end of its sweep it reaches a fidelity no **SAE** attains, even at  $k = 128$ . A few vector-valued concepts carry reconstruction further than many more scalar atoms — the relaxation of the sparsity–reconstruction trade-off the architecture was designed to deliver.

This advantage is specific to the concept axis. Figure 4.2 repeats the comparison in active *neurons*  $\ell_0$  — the decoder-capacity unit (section 4.3) — where the two families do not even occupy the same range. Because each active **VAEE** concept is a vector of width  $d = 64$ , every **VAEE** operating point sits above  $\ell_0 \approx 300$ , while the **SAEs**, with scalar atoms, never exceed  $\ell_0 = 128$ ; a head-to-head reading at matched  $\ell_0$  is not available from these runs. The trend is nonetheless clear — the **SAE** reaches a given reconstruction error with fewer active neurons, since top- $k$  selection targets exactly that scalar count — and the **VAEE** buys its low concept count by activating more neurons per sample. The trade-off is not eliminated but relocated, from the number of concepts a reader inspects to the number of scalars the decoder draws




Figure 4.2: The same frontier in active neurons  $\ell_0$ , the decoder-capacity unit. The families occupy disjoint ranges — SAEs below  $\ell_0 = 128$ , VAEs above  $\ell_0 \approx 300$  — so no matched- $\ell_0$  comparison is available; the SAE reaches a given error with fewer active neurons.

on. Which unit to privilege depends on the purpose: interpretation proceeds concept by concept, never coordinate by coordinate, so  $\kappa$  is the axis on which the comparison bears — and there the VAE wins.

Both metrics presuppose that the active-concept count is well defined — that the continuous inference gate behaves, in practice, as a hard one. Figure 4.3 bears this out: the validation-set distribution of the gate values  $\alpha_k$  is sharply bimodal, almost all of its mass falling within 0.01 of 0 or 1, the extremes one to two orders of magnitude more frequent than any intermediate value. The active-concept count is therefore insensitive to the exact cut-off, and 1/2 is a natural threshold for an activation probability.

#### 4.5.2 QUALITATIVE RESULTS

#### 4.5.3 ABLATION STUDY




Figure 4.3: Validation-set distribution of the  $128 \times 64$  VAE gate values  $\alpha_k$ . The mass is sharply bimodal, concentrated within 0.01 of 0 and 1 (each bin spans 0.01), the extreme bins one to two orders of magnitude more frequent than the interior — so the gate is effectively hard at inference and the active-concept count is insensitive to the  $1/2$  threshold.



## CONCLUSIONS

---

### 5.1 CONTRIBUTIONS

### 5.2 FUTURE WORK





## VAEE DERIVATIONS

---

This appendix collects the full derivations supporting the probabilistic formulation of section 3.3: the optimal form of the gating mechanism (appendix A.1) and the explicit derivation of the evidence lower bound (appendix A.2).

Throughout the appendix,  $\|\cdot\|$  is used instead of  $\|\cdot\|_2$  to avoid making the notation excessively verbose.

### A.1 OPTIMAL GATING MECHANISM

The gate of eq. (3.7) ties a concept’s activation probability to the similarity between its predicted embedding and its prototype. This section shows that this is not an arbitrary heuristic but the *optimal* mask factor of the mean-field posterior (eq. (3.5)), dictated by the Gated Gaussian prior of eq. (3.3).

For a mean-field family, the factor that maximises the ELBO while the remaining factors are held fixed has a closed form: its logarithm equals the expectation of the log-joint over those other factors, up to an additive constant [3]. Applied to the mask factor  $q_\phi(c_k | \mathbf{x})$ , with the expectation taken over the embedding factor  $q_\phi(\mathbf{z}_k | \mathbf{x})$ ,

$$\begin{aligned} \log q_\phi^*(c_k | \mathbf{x}) &= \mathbb{E}_{q_\phi(\mathbf{z}_k | \mathbf{x})} [\log p(c_k, \mathbf{z}_k)] + \text{const} \\ &= \log p(c_k) + \mathbb{E}_{q_\phi(\mathbf{z}_k | \mathbf{x})} [\log p(\mathbf{z}_k | c_k)] + \text{const}. \end{aligned} \quad (\text{A.1})$$

The likelihood  $p_\theta(\mathbf{x} | \mathbf{z})$  of eq. (3.4) carries no  $c_k$  — the generative flow is  $\mathbf{c} \rightarrow \mathbf{z} \rightarrow \mathbf{x}$ , so  $\mathbf{x}$  depends on the mask only through the embeddings — and is therefore absorbed into the constant, leaving only the prior  $p(c_k)$  and the Gated Gaussian term.

Because the mask entry  $c_k$  is binary, the factor  $q_\phi^*(c_k | \mathbf{x})$  is fully characterised by its log-odds: the sigmoid is the inverse of the logit transform, so

$$\begin{aligned} q_\phi^*(c_k = 1 | \mathbf{x}) &= \text{sigmoid} \left( \log \frac{q_\phi^*(c_k = 1 | \mathbf{x})}{1 - q_\phi^*(c_k = 1 | \mathbf{x})} \right) \\ &= \text{sigmoid} \left( \log \frac{q_\phi^*(c_k = 1 | \mathbf{x})}{q_\phi^*(c_k = 0 | \mathbf{x})} \right), \end{aligned}$$

and the optimality condition of eq. (A.1) gives those log-odds as a prior term plus an expected likelihood ratio,

$$\log \frac{q_\phi^*(c_k = 1 | \mathbf{x})}{q_\phi^*(c_k = 0 | \mathbf{x})} = \underbrace{\log \frac{p(c_k = 1)}{p(c_k = 0)}}_{=\text{logit}(\pi)} + \mathbb{E}_{q_\phi(\mathbf{z}_k | \mathbf{x})} \left[ \underbrace{\log \frac{p(\mathbf{z}_k | c_k = 1)}{p(\mathbf{z}_k | c_k = 0)}}_{\text{likelihood ratio}} \right].$$

The prior term is just the logit of the Bernoulli activation rate  $\pi$  of eq. (3.2), since  $\text{logit}(\pi) = \log(\pi/(1-\pi))$ . For the likelihood ratio we substitute the two branches of the Gated Gaussian prior,  $p(\mathbf{z}_k | c_k = 1) = \mathcal{N}(\mathbf{z}_k; \mathbf{h}_k, \sigma_0^2 \mathbb{I})$  and  $p(\mathbf{z}_k | c_k = 0) = \mathcal{N}(\mathbf{z}_k; \mathbf{0}, \sigma_0^2 \mathbb{I})$ . The two share the same isotropic covariance, so their  $(2\pi\sigma_0^2)^{-d/2}$  normalising constants cancel inside the log-ratio, leaving only the exponentials,

$$\begin{aligned} \log \frac{p(\mathbf{z}_k | c_k = 1)}{p(\mathbf{z}_k | c_k = 0)} &= \log \frac{\exp\left(-\frac{1}{2\sigma_0^2} \|\mathbf{z}_k - \mathbf{h}_k\|^2\right)}{\exp\left(-\frac{1}{2\sigma_0^2} \|\mathbf{z}_k\|^2\right)} \\ &= \frac{1}{2\sigma_0^2} \left( \|\mathbf{z}_k\|^2 - \|\mathbf{z}_k - \mathbf{h}_k\|^2 \right). \end{aligned}$$

Expanding the squared distance,  $\|\mathbf{z}_k - \mathbf{h}_k\|^2 = \|\mathbf{z}_k\|^2 - 2\mathbf{z}_k \cdot \mathbf{h}_k + \|\mathbf{h}_k\|^2$ , makes the  $\|\mathbf{z}_k\|^2$  terms cancel exactly, so the likelihood ratio collapses to an inner product against the prototype that is *linear* in the embedding,

$$\log \frac{p(\mathbf{z}_k | c_k = 1)}{p(\mathbf{z}_k | c_k = 0)} = \frac{1}{\sigma_0^2} \left( \mathbf{z}_k \cdot \mathbf{h}_k - \frac{1}{2} \|\mathbf{h}_k\|^2 \right).$$

Linearity is what makes the expectation tractable, and exactly so: with no term beyond first order in  $\mathbf{z}_k$ , the expectation over  $q_\phi(\mathbf{z}_k | \mathbf{x})$  amounts to replacing  $\mathbf{z}_k$  with its mean  $\boldsymbol{\mu}_k(\mathbf{x})$ . Adding the prior log-odds back in and converting through the sigmoid therefore yields the inner-product gate,

$$\alpha_k(\mathbf{x}) = \text{sigmoid} \left( \frac{1}{\sigma_0^2} \left( \boldsymbol{\mu}_k(\mathbf{x}) \cdot \mathbf{h}_k - \frac{1}{2} \|\mathbf{h}_k\|^2 \right) + \text{logit}(\pi) \right). \quad (\text{A.2})$$

To recover the form used in practice (eq. (3.7)), the additive constants  $-\frac{1}{2\sigma_0^2} \|\mathbf{h}_k\|^2 + \text{logit}(\pi)$  are dropped and the remaining scale  $1/\sigma_0^2$  is folded into a single learnable inverse temperature  $1/\tau$ , so that the logit is parameterised as  $(1/\tau) \boldsymbol{\mu}_k(\mathbf{x}) \cdot \mathbf{h}_k$ ; the per-prototype bias is left to be absorbed implicitly into  $\tau$  and into the optimisation of  $\mathbf{h}_k$  itself. Under the exact gate of eq. (A.2) the decision boundary  $\alpha_k = \frac{1}{2}$  sits where  $\boldsymbol{\mu}_k \cdot \mathbf{h}_k = \frac{1}{2} \|\mathbf{h}_k\|^2$ , which is well calibrated only as long as  $\|\boldsymbol{\mu}_k\| \approx \|\mathbf{h}_k\|$ ; otherwise the model can inflate  $\|\boldsymbol{\mu}_k\|$  to push the inner product up and force a concept active. The cosine-similarity variant of eq. (3.11) removes this magnitude coupling by construction, which is why it is the choice of similarity used in this thesis.

A final observation ties this back to the generative model. Since  $x$  depends on the mask only through the embeddings in the generative flow  $c \rightarrow z \rightarrow x$ ,  $c_k$  and  $x$  are conditionally independent given  $z_k$ , so the true posterior factor satisfies  $p(c_k | z_k, x) = p(c_k | z_k)$  — itself a sigmoid of the same linear log-odds in  $z_k$ . The optimal mean-field gate  $\alpha_k(x)$  is thus exactly that posterior factor evaluated at the embedding mean  $z_k = \mu_k(x)$ : conditioning the mask on  $x$  through  $\mu_k$  loses nothing relative to conditioning it on a sampled  $z_k$ , which is why the model gates on the predicted mean directly.

## A.2 EXPLICIT ELBO DERIVATION

This section derives the [ELBO](#) of eq. (3.12) under the mean-field posterior of eq. (3.5). This factorisation is the model’s central variational approximation: the true posterior couples the mask and its embedding through the generative  $c \rightarrow z$  edge, so that even given  $x$  the two remain dependent. Treating them as conditionally independent given  $x$  is what lets all three terms below close in form; it is the only approximation made, every step that follows being exact.

The marginal likelihood requires marginalising over both latents, an integral over the continuous  $z$  and a sum over the  $2^K$  configurations of  $c$  that is intractable. Introducing the approximate posterior and applying Jensen’s inequality<sup>1</sup> to the concave logarithm gives a lower bound:

$$\begin{aligned} \log p_\theta(x) &= \log \sum_c \int p(x, c, z) dz \\ &= \log \sum_c \int p(x, c, z) \frac{q(c, z | x)}{q(c, z | x)} dz \\ &= \log \mathbb{E}_{q(c, z | x)} \left[ \frac{p(x, c, z)}{q(c, z | x)} \right] \\ &\geq \mathbb{E}_{q(c, z | x)} \left[ \log \frac{p(x, c, z)}{q(c, z | x)} \right] \equiv \text{ELBO}. \end{aligned}$$

The generative model factorises over concepts as  $p(x, c, z) = p_\theta(x | z) \prod_k p(c_k) p(z_k | c_k)$ . Substituting this factorisation along with the mean-field posterior and splitting the logarithm of the resulting product gives

$$\begin{aligned} \text{ELBO} &= \mathbb{E}_{q_\phi(c, z | x)} \left[ \log p_\theta(x | z) + \sum_{k=1}^K \log \frac{p(c_k)}{q_\phi(c_k | x)} \right. \\ &\quad \left. + \sum_{k=1}^K \log \frac{p(z_k | c_k)}{q_\phi(z_k | x)} \right]. \end{aligned}$$

<sup>1</sup> Given  $X$  an integrable real-valued random variable and  $\varphi$  a concave function,  $\varphi(\mathbb{E}[X]) \geq \mathbb{E}[\varphi(X)]$ .

Under the mean-field posterior the joint expectation factorises as  $\mathbb{E}_{q_\phi(c, z|x)} = \mathbb{E}_{q_\phi(z|x)} \mathbb{E}_{q_\phi(c|x)}$ , and the three summands separate.

The first carries no mask, so the expectation over  $c$  is trivial and it reduces to the reconstruction term

$$\mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x | z)]. \quad (\text{A.3})$$

The second carries no embedding, so the expectation over  $z$  is trivial and the expectation over  $c$  collects into the negative KL between the Bernoulli posterior and prior,

$$\sum_{k=1}^K \mathbb{E}_{q_\phi(c_k|x)} \left[ \log \frac{p(c_k)}{q_\phi(c_k | x)} \right] = - \sum_{k=1}^K D_{\text{KL}} (q_\phi(c_k | x) \parallel p(c_k)). \quad (\text{A.4})$$

and represents the mask-prior KL. In the third summand the inner expectation over  $z_k$  is itself a KL,

$$\mathbb{E}_{q_\phi(z_k|x)} \left[ \log \frac{p(z_k | c_k)}{q_\phi(z_k | x)} \right] = -D_{\text{KL}} (q_\phi(z_k | x) \parallel p(z_k | c_k)),$$

a function of the still-free  $c_k$ . The outer expectation gives then the conditional-embedding KL term

$$\sum_{k=1}^K \mathbb{E}_{q_\phi(c_k|x)} [D_{\text{KL}} (q_\phi(z_k | x) \parallel p(z_k | c_k))]. \quad (\text{A.5})$$

Collecting the three terms of eqs. (A.3) to (A.5) gives the decomposition of eq. (3.12), reproduced below for convenience. The training objective then follows by negating the ELBO and applying the design choices of section 3.3.2.

$$\begin{aligned} \log p_\theta(x) &\geq \underbrace{\mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x | z)]}_{\text{reconstruction}} + \\ &\quad - \sum_{k=1}^K \underbrace{D_{\text{KL}} (q_\phi(c_k | x) \parallel p(c_k))}_{\text{mask-prior KL}} + \\ &\quad - \sum_{k=1}^K \underbrace{\mathbb{E}_{q_\phi(c_k|x)} [D_{\text{KL}} (q_\phi(z_k | x) \parallel p(z_k | c_k))]}_{\text{conditional-embedding KL}}. \end{aligned} \quad (\text{3.12 revisited})$$

**CLOSED-FORM CONDITIONAL KL** In eq. (A.5) the expectation over the Bernoulli  $\mathbb{E}_{q_\phi(c_k|x)}$  is the convex combination of its two branches with weights  $\alpha_k$  and  $1 - \alpha_k$ :

$$\begin{aligned} \sum_{k=1}^K \mathbb{E}_{q_\phi(c_k|x)} [D_{\text{KL}} (q_\phi(z_k | x) \parallel p(z_k | c_k))] &= \\ &= \sum_{k=1}^K \left[ \alpha_k D_{\text{KL}}^{(k,1)} + (1 - \alpha_k) D_{\text{KL}}^{(k,0)} \right], \end{aligned}$$

where  $D_{\text{KL}}^{(k,c)} = D_{\text{KL}}(q_\phi(\mathbf{z}_k | \mathbf{x}) \| p(\mathbf{z}_k | c_k=c))$ . Both branches are Kullback–Leiblers between Gaussians sharing the covariance  $\sigma_0^2 \mathbb{I}$ , for which the divergence reduces to the scaled squared distance between the means,

$$D_{\text{KL}}(\mathcal{N}(\boldsymbol{\mu}, \sigma_0^2 \mathbb{I}) \| \mathcal{N}(\boldsymbol{\nu}, \sigma_0^2 \mathbb{I})) = \frac{1}{2\sigma_0^2} \|\boldsymbol{\mu} - \boldsymbol{\nu}\|^2$$

With the prior means  $c_k \mathbf{h}_k$  of eq. (3.3) — the prototype  $\mathbf{h}_k$  when active, the origin when inactive — the conditional-embedding [KL](#) takes its closed form,

$$\begin{aligned} \sum_{k=1}^K \mathbb{E}_{q_\phi(c_k|\mathbf{x})} [D_{\text{KL}}(q_\phi(\mathbf{z}_k | \mathbf{x}) \| p(\mathbf{z}_k | c_k))] &= \\ &= \frac{1}{2\sigma_0^2} \sum_{k=1}^K [\alpha_k \|\boldsymbol{\mu}_k - \mathbf{h}_k\|^2 + (1 - \alpha_k) \|\boldsymbol{\mu}_k\|^2], \end{aligned}$$

which is the term substituted into the practical loss of eq. (3.13), with the factor  $\frac{1}{2\sigma_0^2}$  being absorbed into the  $\gamma$  hyperparameter.



## BIBLIOGRAPHY

---

- [1] Emmanuel Ameisen et al. *Circuit Tracing: Revealing Computational Graphs in Language Models*. Tech. rep. Transformer Circuits Thread, Mar. 2025. URL: <https://transformer-circuits.pub/2025/attribution-graphs/methods.html> (visited on 11/06/2026).
- [2] Steven Bills, Nick Cammarata, Dan Mossing, Henk Tillman, Leo Gao, Gabriel Goh, Ilya Sutskever, Jan Leike, Jeff Wu and William Saunders. *Language Models Can Explain Neurons in Language Models*. Tech. rep. OpenAI, May 2023. URL: <https://openaipublic.blob.core.windows.net/neuron-explainer/paper/index.html> (visited on 11/06/2026).
- [3] David M. Blei, Alp Kucukelbir and Jon D. McAuliffe. ‘Variational Inference: A Review for Statisticians’. In: *Journal of the American Statistical Association* 112.518 (Apr. 2017), pp. 859–877. ISSN: 0162-1459. DOI: [10.1080/01621459.2017.1285773](https://doi.org/10.1080/01621459.2017.1285773). URL: <https://doi.org/10.1080/01621459.2017.1285773> (visited on 17/06/2026).
- [4] Trenton Bricken et al. *Towards Monosemanticity: Decomposing Language Models With Dictionary Learning*. Tech. rep. Transformer Circuits Thread, Oct. 2023. URL: <https://transformer-circuits.pub/2023/monosemantic-features/index.html> (visited on 11/06/2026).
- [5] Tom Brown et al. ‘Language Models Are Few-Shot Learners’. In: *Advances in Neural Information Processing Systems*. Vol. 33. Curran Associates, Inc., 2020, pp. 1877–1901. URL: <https://proceedings.neurips.cc/paper/2020/hash/1457c0d6bfc4967418bfb8ac142f64a-Abstract.html> (visited on 13/06/2026).
- [6] Francesco De Santis, Philippe Bich, Gabriele Ciravegna, Pietro Barbiero, Tania Cerquitelli and Danilo Giordano. ‘Towards Better Generalization and Interpretability in Unsupervised Concept-Based Models’. In: *Machine Learning and Knowledge Discovery in Databases. Research Track: European Conference, ECML PKDD 2025, Porto, Portugal, September 15–19, 2025, Proceedings, Part III*. Berlin, Heidelberg: Springer-Verlag, Oct. 2025, pp. 478–494. ISBN: 978-3-032-06065-5. DOI: [10.1007/978-3-032-06066-2\\_28](https://doi.org/10.1007/978-3-032-06066-2_28). URL: [https://doi.org/10.1007/978-3-032-06066-2\\_28](https://doi.org/10.1007/978-3-032-06066-2_28) (visited on 11/06/2026).

- [7] Qingxiu Dong et al. ‘A Survey on In-context Learning’. In: *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Ed. by Yaser Al-Onaizan, Mohit Bansal and Yun-Nung Chen. Miami, Florida, USA: Association for Computational Linguistics, Nov. 2024, pp. 1107–1128. DOI: [10.18653/v1/2024.emnlp-main.64](https://doi.org/10.18653/v1/2024.emnlp-main.64). URL: <https://aclanthology.org/2024.emnlp-main.64/> (visited on 11/06/2026).
- [8] Jacob Dunefsky, Philippe Chlenski and Neel Nanda. ‘Transcoders Find Interpretable LLM Feature Circuits’. In: *Advances in Neural Information Processing Systems* 37 (Dec. 2024), pp. 24375–24410. DOI: [10.52202/079017-0768](https://doi.org/10.52202/079017-0768). URL: [https://proceedings.neurips.cc/paper\\_files/paper/2024/hash/2b8f4db0464cc5b6e9d5e6bea4b9f308-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2024/hash/2b8f4db0464cc5b6e9d5e6bea4b9f308-Abstract-Conference.html) (visited on 10/06/2026).
- [9] Nelson Elhage et al. *A Mathematical Framework for Transformer Circuits*. Tech. rep. Transformer Circuits Thread, Dec. 2021. URL: <https://transformer-circuits.pub/2021/framework/index.html> (visited on 13/06/2026).
- [10] Nelson Elhage et al. *Toy Models of Superposition*. Tech. rep. Transformer Circuits Thread, Sept. 2022. URL: [https://transformer-circuits.pub/2022/toy\\_model/index.html](https://transformer-circuits.pub/2022/toy_model/index.html) (visited on 11/06/2026).
- [11] Mateo Espinosa Zarlenga et al. ‘Concept Embedding Models: Beyond the Accuracy-Explainability Trade-Off’. In: *Advances in Neural Information Processing Systems* 35 (Dec. 2022), pp. 21400–21413. URL: [https://proceedings.neurips.cc/paper\\_files/paper/2022/hash/867c06823281e506e8059f5c13a57f75-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2022/hash/867c06823281e506e8059f5c13a57f75-Abstract-Conference.html) (visited on 12/06/2026).
- [12] Leo Gao, Tom Dupre la Tour, Henk Tillman, Gabriel Goh, Rajan Troll, Alec Radford, Ilya Sutskever, Jan Leike and Jeffrey Wu. ‘Scaling and Evaluating Sparse Autoencoders’. In: *The Thirteenth International Conference on Learning Representations*. Oct. 2024. URL: <https://openreview.net/forum?id=tcsZt9ZNKD> (visited on 14/05/2026).
- [13] Ian Goodfellow, Aaron Courville and Yoshua Bengio. *Deep Learning*. Adaptive Computation and Machine Learning. Cambridge, Massachusetts: The MIT Press, 2016. ISBN: 978-0-262-03561-3 978-0-262-33737-3.
- [14] Irina Higgins, Loic Matthey, Arka Pal, Christopher Burgess, Xavier Glorot, Matthew Botvinick, Shakir Mohamed and Alexander Lerchner. ‘Beta-VAE: Learning Basic Visual Concepts with a Constrained Variational Framework’. In: *International Conference on Learning Representations*. Feb. 2017. URL: <https://openreview.net/forum?id=Sy2fzU9gl> (visited on 10/06/2026).



- [15] Robert Huben, Hoagy Cunningham, Logan Riggs Smith, Aidan Ewart and Lee Sharkey. ‘Sparse Autoencoders Find Highly Interpretable Features in Language Models’. In: *The Twelfth International Conference on Learning Representations*. Oct. 2023. URL: <https://openreview.net/forum?id=F76bwRSLeK> (visited on 28/05/2026).
- [16] Eric Jang, Shixiang Gu and Ben Poole. ‘Categorical Reparameterization with Gumbel-Softmax’. In: *International Conference on Learning Representations*. Feb. 2017. URL: <https://openreview.net/forum?id=rkE3y85ee> (visited on 14/05/2026).
- [17] Albert Q. Jiang et al. *Mistral 7B*. Oct. 2023. DOI: [10.48550/arXiv.2310.06825](https://doi.org/10.48550/arXiv.2310.06825). arXiv: [2310.06825](https://arxiv.org/abs/2310.06825) [cs.CL]. URL: <http://arxiv.org/abs/2310.06825> (visited on 18/06/2026).
- [18] Been Kim, Rajiv Khanna and Oluwasanmi O Koyejo. ‘Examples Are Not Enough, Learn to Criticize! Criticism for Interpretability’. In: *Advances in Neural Information Processing Systems*. Vol. 29. Curran Associates, Inc., 2016. URL: [https://proceedings.neurips.cc/paper\\_files/paper/2016/hash/5680522b8e2bb01943234bce7bf84534-Abstract.html](https://proceedings.neurips.cc/paper_files/paper/2016/hash/5680522b8e2bb01943234bce7bf84534-Abstract.html) (visited on 11/06/2026).
- [19] Been Kim, Martin Wattenberg, Justin Gilmer, Carrie Cai, James Wexler, Fernanda Viegas and Rory Sayres. ‘Interpretability Beyond Feature Attribution: Quantitative Testing with Concept Activation Vectors (TCAV)’. In: *Proceedings of the 35th International Conference on Machine Learning*. PMLR, July 2018, pp. 2668–2677. URL: <https://proceedings.mlr.press/v80/kim18d.html> (visited on 11/06/2026).
- [20] Diederik P. Kingma and Jimmy Ba. *Adam: A Method for Stochastic Optimization*. Jan. 2017. DOI: [10.48550/arXiv.1412.6980](https://doi.org/10.48550/arXiv.1412.6980). arXiv: [1412.6980](https://arxiv.org/abs/1412.6980) [cs.LG]. URL: <http://arxiv.org/abs/1412.6980> (visited on 19/06/2026).
- [21] Diederik P. Kingma and Max Welling. ‘Auto-Encoding Variational Bayes’. In: *International Conference on Learning Representations*, May 2014. DOI: [10.48550/arXiv.1312.6114](https://doi.org/10.48550/arXiv.1312.6114). arXiv: [1312.6114](https://arxiv.org/abs/1312.6114) [stat]. URL: <http://arxiv.org/abs/1312.6114> (visited on 07/11/2025).
- [22] Pang Wei Koh, Thao Nguyen, Yew Siang Tang, Stephen Mussmann, Emma Pierson, Been Kim and Percy Liang. ‘Concept Bottleneck Models’. In: *Proceedings of the 37th International Conference on Machine Learning*. PMLR, Nov. 2020, pp. 5338–5348. URL: <https://proceedings.mlr.press/v119/koh20a.html> (visited on 11/06/2026).

- [23] Scott M Lundberg and Su-In Lee. ‘A Unified Approach to Interpreting Model Predictions’. In: *Advances in Neural Information Processing Systems*. Vol. 30. Curran Associates, Inc., 2017. URL: [https://proceedings.neurips.cc/paper\\_files/paper/2017/hash/8a20a8621978632d76c43dfd28b67767-Abstract.html](https://proceedings.neurips.cc/paper_files/paper/2017/hash/8a20a8621978632d76c43dfd28b67767-Abstract.html) (visited on 01/06/2026).
- [24] Chris J. Maddison, Andriy Mnih and Yee Whye Teh. ‘The Concrete Distribution: A Continuous Relaxation of Discrete Random Variables’. In: *International Conference on Learning Representations*. Feb. 2017. URL: <https://openreview.net/forum?id=S1jE5L5gl> (visited on 14/05/2026).
- [25] Samuel Marks, Can Rager, Eric Michaud, Yonatan Belinkov, David Bau and Aaron Mueller. ‘Sparse Feature Circuits: Discovering and Editing Interpretable Causal Graphs in Language Models’. In: *International Conference on Learning Representations 2025* (May 2025), pp. 23888–23923. URL: [https://proceedings.iclr.cc/paper\\_files/paper/2025/hash/3ba4d47a83e498c2b1a0868cba20f6de-Abstract-Conference.html](https://proceedings.iclr.cc/paper_files/paper/2025/hash/3ba4d47a83e498c2b1a0868cba20f6de-Abstract-Conference.html) (visited on 10/06/2026).
- [26] Christoph Molnar. *Interpretable Machine Learning: A Guide for Making Black Box Models Explainable*. 3rd edition. Munich, Germany: Christoph Molnar, 2025. ISBN: 978-3-911578-03-5.
- [27] Tuomas Oikarinen, Subhro Das, Lam M. Nguyen and Tsui-Wei Weng. ‘Label-Free Concept Bottleneck Models’. In: *The Eleventh International Conference on Learning Representations*. Sept. 2022. URL: <https://openreview.net/forum?id=FlCg47MNvBA> (visited on 11/06/2026).
- [28] Chris Olah, Nick Cammarata, Ludwig Schubert, Gabriel Goh, Michael Petrov and Shan Carter. ‘Zoom In: An Introduction to Circuits’. In: *Distill* 5.3 (Mar. 2020), e00024.001. ISSN: 2476-0757. DOI: [10.23915/distill.00024.001](https://doi.org/10.23915/distill.00024.001). URL: <https://distill.pub/2020/circuits/zoom-in> (visited on 04/05/2026).
- [29] Bruno A. Olshausen and David J. Field. ‘Sparse Coding with an Overcomplete Basis Set: A Strategy Employed by V1?’ In: *Vision Research* 37.23 (Dec. 1997), pp. 3311–3325. ISSN: 0042-6989. DOI: [10.1016/S0042-6989\(97\)00169-7](https://doi.org/10.1016/S0042-6989(97)00169-7). URL: <https://www.sciencedirect.com/science/article/pii/S0042698997001697> (visited on 28/05/2026).
- [30] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei and Ilya Sutskever. *Language Models Are Unsupervised Multitask Learners*. Tech. rep. OpenAI, 2019. (Visited on 13/06/2026).

- [31] Senthooran Rajamanoharan, Arthur Conmy, Lewis Smith, Tom Lieberum, Vikrant Varma, János Kramár, Rohin Shah and Neel Nanda. ‘Improving Sparse Decomposition of Language Model Activations with Gated Sparse Autoencoders’. In: *Advances in Neural Information Processing Systems* 37 (Dec. 2024), pp. 775–818. DOI: [10.52202/079017-0024](https://doi.org/10.52202/079017-0024). URL: [https://proceedings.neurips.cc/paper\\_files/paper/2024/hash/01772a8b0420baec00c4d59fe2fbace6-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2024/hash/01772a8b0420baec00c4d59fe2fbace6-Abstract-Conference.html) (visited on 05/06/2026).
- [32] Senthooran Rajamanoharan, Tom Lieberum, Nicolas Sonnerat, Arthur Conmy, Vikrant Varma, János Kramár and Neel Nanda. *Jumping Ahead: Improving Reconstruction Fidelity with JumpReLU Sparse Autoencoders*. Aug. 2024. DOI: [10.48550/arXiv.2407.14435](https://doi.org/10.48550/arXiv.2407.14435). arXiv: [2407.14435 \[cs\]](https://arxiv.org/abs/2407.14435). URL: <http://arxiv.org/abs/2407.14435> (visited on 23/02/2026).
- [33] Danilo Jimenez Rezende, Shakir Mohamed and Daan Wierstra. ‘Stochastic Backpropagation and Approximate Inference in Deep Generative Models’. In: *Proceedings of the 31st International Conference on Machine Learning*. PMLR, June 2014, pp. 1278–1286. URL: <https://proceedings.mlr.press/v32/rezende14.html> (visited on 10/06/2026).
- [34] Marco Tulio Ribeiro, Sameer Singh and Carlos Guestrin. “‘Why Should I Trust You?’: Explaining the Predictions of Any Classifier”. In: *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. KDD ’16. New York, NY, USA: Association for Computing Machinery, Aug. 2016, pp. 1135–1144. ISBN: 978-1-4503-4232-2. DOI: [10.1145/2939672.2939778](https://doi.org/10.1145/2939672.2939778). URL: <https://dl.acm.org/doi/10.1145/2939672.2939778> (visited on 01/06/2026).
- [35] Cynthia Rudin. ‘Stop Explaining Black Box Machine Learning Models for High Stakes Decisions and Use Interpretable Models Instead’. In: *Nature Machine Intelligence* 1.5 (May 2019), pp. 206–215. ISSN: 2522-5839. DOI: [10.1038/s42256-019-0048-x](https://doi.org/10.1038/s42256-019-0048-x). URL: <https://www.nature.com/articles/s42256-019-0048-x> (visited on 01/06/2026).
- [36] Ramprasaath R. Selvaraju, Michael Cogswell, Abhishek Das, Ramakrishna Vedantam, Devi Parikh and Dhruv Batra. ‘Grad-CAM: Visual Explanations from Deep Networks via Gradient-Based Localization’. In: *2017 IEEE International Conference on Computer Vision (ICCV)*. Oct. 2017, pp. 618–626. DOI: [10.1109/ICCV.2017.74](https://doi.org/10.1109/ICCV.2017.74). URL: <https://ieeexplore.ieee.org/document/8237336> (visited on 01/06/2026).
- [37] Irhum Shafkat. *Intuitively Understanding Variational Autoencoders*. Feb. 2018. URL: <https://medium.com/data-science/>

- intuitively - understanding - variational - autoencoders - 1bfe67eb5daf (visited on 14/05/2026).
- [38] Karen Simonyan, Andrea Vedaldi and Andrew Zisserman. *Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps*. Apr. 2014. DOI: [10.48550/arXiv.1312.6034](https://doi.org/10.48550/arXiv.1312.6034). arXiv: [1312.6034](https://arxiv.org/abs/1312.6034) [cs.CV]. URL: <http://arxiv.org/abs/1312.6034> (visited on 11/06/2026).
  - [39] Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D. Manning, Andrew Ng and Christopher Potts. 'Recursive Deep Models for Semantic Compositionality Over a Sentiment Treebank'. In: *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*. Ed. by David Yarowsky, Timothy Baldwin, Anna Korhonen, Karen Livescu and Steven Bethard. Seattle, Washington, USA: Association for Computational Linguistics, Oct. 2013, pp. 1631–1642. URL: <https://aclanthology.org/D13-1170/> (visited on 18/06/2026).
  - [40] Chung-En Sun, Tuomas Oikarinen, Berk Ustun and Tsui-Wei Weng. 'Concept Bottleneck Large Language Models'. In: *International Conference on Learning Representations 2025* (May 2025), pp. 89371–89411. URL: [https://proceedings.iclr.cc/paper\\_files/paper/2025/hash/de4ce91dfe56b919ee1c228d6a78f866-Abstract-Conference.html](https://proceedings.iclr.cc/paper_files/paper/2025/hash/de4ce91dfe56b919ee1c228d6a78f866-Abstract-Conference.html) (visited on 01/06/2026).
  - [41] Adly Templeton et al. *Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet*. URL: <https://transformer-circuits.pub/2024/scaling-monosemanticity/index.html#appendix-citation> (visited on 20/10/2025).
  - [42] Robert Tibshirani. 'Regression Shrinkage and Selection Via the Lasso'. In: *Journal of the Royal Statistical Society: Series B (Methodological)* 58.1 (Jan. 1996), pp. 267–288. ISSN: 0035-9246. DOI: [10.1111/j.2517-6161.1996.tb02080.x](https://doi.org/10.1111/j.2517-6161.1996.tb02080.x). URL: <https://doi.org/10.1111/j.2517-6161.1996.tb02080.x> (visited on 29/05/2026).
  - [43] Aaron van den Oord, Oriol Vinyals and Koray Kavukcuoglu. 'Neural Discrete Representation Learning'. In: *Advances in Neural Information Processing Systems*. Vol. 30. Curran Associates, Inc., 2017. URL: <https://papers.nips.cc/paper/2017/hash/7a98af17e63a0ac09ce2e96d03992fbc-Abstract.html> (visited on 11/06/2026).
  - [44] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser and Illia Polosukhin. 'Attention Is All You Need'. In: *Advances in Neural Information Processing Systems*. Vol. 30. Curran Associates, Inc., 2017. URL: [https://proceedings.neurips.cc/paper\\_files/paper/2017/hash/3f5ee243547dee91fbd053c1c4a845aa-Abstract.html](https://proceedings.neurips.cc/paper_files/paper/2017/hash/3f5ee243547dee91fbd053c1c4a845aa-Abstract.html) (visited on 11/06/2026).

- [45] Sandra Wachter, Brent Mittelstadt and Chris Russell. *Counterfactual Explanations Without Opening the Black Box: Automated Decisions and the GDPR*. SSRN Scholarly Paper. Rochester, NY, Oct. 2017. DOI: [10.2139/ssrn.3063289](https://doi.org/10.2139/ssrn.3063289). Social Science Research Network: [3063289](https://papers.ssrn.com/abstract=3063289). URL: <https://papers.ssrn.com/abstract=3063289> (visited on 11/06/2026).
- [46] Benjamin Wright and Lee Sharkey. ‘Addressing Feature Suppression in SAEs — AI Alignment Forum’. In: (Feb. 2024). URL: <https://www.alignmentforum.org/posts/3JuSjTZyMzaSeTxKk/addressing-feature-suppression-in-saes> (visited on 29/05/2026).



## COLOPHON

This document was typeset using the typographical look-and-feel `classicthesis` developed by André Miede and Ivo Pletikosić. The style was inspired by Robert Bringhurst's seminal book on typography "*The Elements of Typographic Style*". `classicthesis` is available for both L<sup>A</sup>T<sub>E</sub>X and L<sup>y</sup>X:

<https://bitbucket.org/amiede/classicthesis/>

Happy users of `classicthesis` usually send a real postcard to the author, a collection of postcards received so far is featured here:

<http://postcards.miede.de/>

Thank you very much for your feedback and contribution.