



**UNIVERSITÀ
DI TORINO**

UNIVERSITÀ DEGLI STUDI DI TORINO

Corso di Laurea Magistrale in Fisica dei Sistemi Complessi

Improving VAE Representational Capabilities for Understanding LLMs

Tesi di Laurea Magistrale

Relatore:

Prof. Michele Caselle

Corelatori:

Prof. Alan Perotti

Prof. André Panisson

Controrelatore:

Prof. Matteo Osella

Candidato:

Francesco Marchisotti

Tutor Aziendale:

Gabriele Ciravegna

Anno Accademico 2025/2026

Francesco Marchisotti: *Improving VAE Representational Capabilities for Understanding LLMs*, Fisica dei Sistemi Complessi

CENTAI TEAM:

Gabriele Ciravegna

Alan Perotti

André Panisson

Francesco Cozzi

UNITO PROFESSORS:

Michele Caselle

Matteo Osella

LOCATION:

Torino

TIME FRAME:

Oct. 2025 - Jul. 2026

“I wish there was a way to know you’re in the good old days
before you’ve actually left them.”

— *Andy Bernard*

ABSTRACT

Concept-based methods represent a model’s computation in terms of human-meaningful concepts, but they typically encode each concept as a single scalar coefficient, e. g., a unit in the bottleneck of a Concept Bottleneck Model, or an atom in the dictionary of a Sparse AutoEncoder (SAE). This scalar encoding trades representational capacity for simplicity: each concept can vary only in magnitude, so reconstructing the underlying representation faithfully tends to require many concepts active at once. This thesis introduces the Variational Auto Embedding Encoder (VAEE), a model derived from first principles that represents each concept as a learned vector embedding instead, drawing together ideas from Concept Embedding Models, Variational AutoEncoders, and Vector-Quantized Variational AutoEncoders. The encoder maps an input to one vector embedding per concept; a learned prototype for each concept then decides, by similarity, which concepts are active, and the input is reconstructed from the stack of active embeddings. Its capabilities are demonstrated in one application, the decomposition of Large Language Model activations, where a VAEE with 8 active concepts matches the reconstruction of a parameter-matched TopK SAE that requires 128 — evidence that vector-valued concepts carry more information per unit of sparsity. The thesis is deliberate about the limits of this evidence: reconstruction was assessed only by mean squared error, each concept now spans several active neurons rather than one, and the recovered concepts were not found more interpretable than the SAE’s. Its contribution, then, lies in this added representational capacity; the interpretability of its concepts, their use for steering, and its application to representation learning more broadly are left for future work.

ACKNOWLEDGMENTS

To whom shall be acknowledged,
you are being acknowledged with this exact statement.

“I’m so funny.”

— *The candidate, who’s not actually funny*

CONTENTS

1	INTRODUCTION	1
1.1	Motivation	1
1.2	Thesis Structure	1
2	BACKGROUND AND RELATED WORK	3
2.1	Natural Language Processing	3
2.1.1	Language Modelling	3
2.1.2	The Transformer	4
2.2	eXplainable AI	6
2.2.1	Traditional XAI	6
2.2.2	Concept-based XAI	7
2.3	AutoEncoders	10
2.3.1	Standard AutoEncoders	10
2.3.2	Variational AutoEncoders	12
2.3.3	Sparse AutoEncoders	15
3	METHODS	21
3.1	Aim of the Work	21
3.2	Definitions/Notation	22
3.3	Model Development	22
3.3.1	Probabilistic Formulation	22
3.3.2	Architectures	22
4	EXPERIMENTS	23
4.1	Model Training	23
4.2	Evaluation Metrics	23
4.2.1	MSE	23
4.2.2	ℓ_0	23
4.2.3	Perplexity	23
4.3	Results	23
5	CONCLUSIONS	25
5.1	Contributions	25
5.2	Future Work	25
A	VAEE DERIVATIONS	27
A.1	Optimal Gating Mechanism	27
A.2	Explicit ELBO Derivation	27
	BIBLIOGRAPHY	29

LIST OF FIGURES

Figure 2.1	A decoder-only transformer drawn around the residual stream	5
Figure 2.2	Denoising autoencoder architecture	11
Figure 2.3	Autoencoder and VAE latent spaces on MNIST	12
Figure 2.4	A polysemantic neuron in InceptionV1	16

LIST OF TABLES

LISTINGS

ACRONYMS

CB-LLM	Concept Bottleneck LLM	8 f.
CBM	Concept Bottleneck Model	v , 8, 20 f.
CE	Cross-Entropy	3, 20
CEM	Concept Embedding Model	v , 8, 20 ff.
ELBO	Evidence Lower Bound	13, 15
KL	Kullback–Leibler	13 ff.
LCBM	Learnable Concept-based Model	21
LF-CBM	Label-free CBM	8
LLM	Large Language Model	v , 3 f., 6, 8 f., 19, 21 f.

MECHINT Mechanistic Interpretability	9, 16, 19
MLP Multi-Layer Perceptron	4, 19
MSE mean squared error	v, 10, 13
PCA Principal Component Analysis	10
SAE Sparse AutoEncoder	v, 9, 15–21
VAE Variational AutoEncoder	v, 12–15, 18, 20, 22
VAEE Variational Auto Embedding Encoder	v, 22
VQ-VAE Vector-Quantized Variational AutoEncoder .	v, 14, 20, 22

INTRODUCTION

1.1 MOTIVATION

1.2 THESIS STRUCTURE

BACKGROUND AND RELATED WORK

This chapter establishes the background on which the thesis builds and reviews the related work, moving from the setting in which the contribution is demonstrated to the ideas from which it is built. Section 2.1 introduces Large Language Models and the transformer architecture that produces the dense internal activations on which that demonstration rests. Section 2.2 reviews explainable AI, tracing the field’s shift from attributing a prediction to input features towards explaining it in terms of human-meaningful *concepts*, the view this work adopts. Section 2.3 then develops the autoencoder family, from the standard autoencoder through the variational and sparse variants between which the thesis’s contribution sits. A single tension runs through all three: reconstructing a model’s activations faithfully while decomposing them into a small, interpretable set of concepts — the tension the method of section 3.3 sets out to relax.

2.1 NATURAL LANGUAGE PROCESSING

Natural Language Processing is the branch of machine learning concerned with getting computers to process and generate human language. For most of its history the field advanced through task-specific systems — separate models for translation, sentiment analysis, or question answering — but it has since converged on a single recipe: train one large neural network to predict text on a vast corpus, then reuse it across tasks. The Large Language Models (LLMs) that result are the systems on which the thesis’s contribution is demonstrated. What follows introduces only as much of their training and architecture as the later chapters rely on, with the emphasis on the internal activations these models compute rather than the text they emit.

2.1.1 LANGUAGE MODELLING

A language model defines a probability distribution over sequences of text. In the dominant *autoregressive* formulation, this distribution is factorised left to right, so that the model repeatedly predicts the next unit of text given all the units before it [33]. The units are *tokens*, words or sub-word fragments drawn from a fixed vocabulary, and the model is a neural network that maps a sequence of tokens to a probability distribution over the vocabulary for the next position. Training minimises the Cross-Entropy (CE) between this predicted

distribution and the true next token, averaged over a large corpus; the objective requires no human annotation, since the supervision signal is the text itself.

Scaling this recipe to more parameters, more data, and more computing power yields LLMs whose capabilities extend well beyond the literal next-token task they were trained on. A network trained only to continue text acquires the ability to translate, summarise, answer questions, and follow instructions, often from nothing more than a description and a few examples supplied in its input, with no change to its weights [7, 9]. This generality is what makes LLMs both powerful and hard to interpret: a single fixed network comes to encode a vast range of linguistic and world knowledge, distributed across its parameters and activations in a form that no direct reading of the weights makes legible.

2.1.2 THE TRANSFORMER

Modern LLMs are almost universally built on the *transformer* architecture [49]. It embeds the input tokens into a sequence of vectors, one per position, and refines them through a stack of identical *blocks* before reading the final vectors out as next-token distributions. Each block is made of two sublayers with complementary roles: an *attention* sublayer, which exchanges information between positions, and a Multi-Layer Perceptron (MLP) sublayer, which transforms each position on its own. Figure 2.1 shows this organisation.

ATTENTION *Attention* is the mechanism by which positions exchange information. At each position, an *attention head* forms a query and compares it against a key emitted by every position in the sequence; the better a position matches, the more it contributes to a weighted average of values gathered from across the sequence, and that average becomes the head’s output for the current position [49]. A sublayer runs several such heads in parallel, each free to specialise (one might link a pronoun to its referent, another a verb to its subject). Because attention is the only component that moves information between positions, it is where the model assembles context: it turns a position’s representation from a fact about a single token into one that reflects the surrounding text.

MLP SUBLAYERS Where attention mixes information across positions, an MLP sublayer processes each position on its own, applying the same feed-forward transformation — a projection up to a wider hidden layer, a nonlinearity, and a projection back down — to the vector at every position independently. Much of the model’s capacity sits in these sublayers, which are widely thought to store and retrieve




Figure 2.1: A decoder-only transformer, drawn around the *residual stream*. Each token is embedded into a vector; the resulting per-position state flows up the central stream, and every block reads from the stream and adds its result back through a residual connection — first a multi-head self-attention sublayer, which mixes information across positions, then a position-wise **MLP** sublayer. After N such blocks the final stream state is unembedded into a next-token distribution. The additive structure of eq. (2.1) is what makes the residual stream a single, well-defined activation space to decompose; the view follows Elhage et al. [11].

its learned associations and are a primary site at which interpretable features appear [6].

THE RESIDUAL STREAM A single detail ties these sublayers together. Neither sublayer overwrites the vector it is given; each computes an update and *adds* it back, through a *residual connection*. Every position therefore carries one running vector — the *residual stream* — that every sublayer of every block reads from and writes to. Let $\mathbf{x}_t^{(\ell)}$ denote the residual stream at position t as it enters block ℓ ; the block updates it as

$$\mathbf{x}_t^{(\ell+1)} = \mathbf{x}_t^{(\ell)} + \text{Attn}^{(\ell)} \left(\mathbf{x}_t^{(\ell)}; \left\{ \mathbf{x}_{t'}^{(\ell)} \right\}_{t'=1, \dots, t-1, t+1, \dots, T} \right) + \text{MLP}^{(\ell)} \left(\mathbf{x}_t^{(\ell)} \right), \quad (2.1)$$

where $\text{Attn}^{(\ell)}$ also reads the residual stream at the other positions. The stream starts as the token’s embedding and accumulates the output of every sublayer in turn; because attention writes into it, a position’s stream is never a fact about that token alone but a running summary of the context around it. This additive structure is what makes

the residual stream a single, well-defined space in which to study the model’s computation — the view developed in the mechanistic-interpretability literature [11] and adopted throughout this thesis. It is also the object the later chapters take apart: the activation vector $\mathbf{x} \in \mathbb{R}^n$ that the autoencoders of section 2.3 receive as input is exactly such a residual-stream state $\mathbf{x}^{(\ell)}$, read off at a chosen layer ℓ .

2.2 EXPLAINABLE AI

The incredible performance of modern machine-learning systems comes at the cost of transparency: a deep network’s output is the product of millions of interacting parameters, with no accompanying account of why it was produced. eXplainable Artificial Intelligence (XAI) is the field that sets out to supply such accounts, making a model’s behaviour intelligible enough to be trusted, debugged, or learned from. The methods relevant to this thesis are most naturally organised around a shift in what counts as an explanation: from *attribution*, which asks *where* (“which input features drove a prediction?”), to *concepts*, the human-meaningful units in terms of which a model can be said to reason. It is this latter, concept-based view on which this work builds.

2.2.1 TRADITIONAL XAI

2.2.1.1 *Intrinsically Interpretable Models*

The most direct route to an interpretable model is to choose one whose decision process is transparent by construction. Linear and logistic regression, decision trees, naive Bayes classifiers, and k -nearest neighbours are the canonical examples: their parameters or structure can be inspected directly, so that the explanation *is* the model rather than a separate artefact derived from it [29]. Rudin [39] argues that such models should be preferred wherever they are competitive, since post-hoc explanations of an opaque model are approximations that need not be faithful to its actual computation. This route is, however, closed for the systems studied in this thesis: LLMs comprise billions of nonlinearly interacting parameters, and no direct reading of their weights yields a human-understandable account of their behaviour. Interpretability must therefore be recovered *post hoc*, from a model that is already trained and fixed.

2.2.1.2 *Post-hoc Attribution Methods*

The dominant post-hoc paradigm explains an individual prediction by *attributing* it to the input features, assigning each feature a score that quantifies its contribution to the output. LIME [37] does this

by fitting a simple, interpretable surrogate (typically a sparse linear model) to the behaviour of the black box in a local neighbourhood of the instance being explained. SHAP [27] unifies this and several related schemes under the Shapley value from cooperative game theory, yielding feature attributions that are additive and satisfy a set of desirable axioms. For models operating on images, gradient-based saliency maps [43] and their refinements such as Grad-CAM [41] produce a heatmap over the input that highlights the regions most influential to the prediction. What these methods share is also their central limitation: they indicate *where* in the input the model is sensitive, but not *what* the model has internally come to represent. An attribution map can mark a pixel or a token as important without revealing the human-meaningful concept it participates in: a gap that motivates the concept-based methods of section 2.2.2.

COUNTERFACTUAL EXPLANATIONS A counterfactual explanation answers a contrastive question: what is the smallest change to the input that would have produced a different, desired prediction [51]? Rather than scoring the features of the present input, it points to a nearby instance on the other side of the decision boundary: “had this applicant’s income been slightly higher, the loan would have been granted”. Such explanations are attractive because they are actionable and require no access to the model’s internals, only the ability to query it. Generating counterfactuals that are at the same time close to the original input and realistic is itself a modelling problem, one which generative models are naturally suited to solve.

2.2.1.3 Prototypes

A complementary, example-based strategy explains a model not through feature scores but through representative instances. A *prototype* is a data point that is typical of a region of the input space, chosen so that a small set of such points summarises the data (or the model’s behaviour upon it) in terms a human can grasp by inspection [21]. The defining characteristic, for the purposes of this thesis, is that a prototype in this traditional sense is an *input instance*: an actual or representative example drawn from the data space.

2.2.2 CONCEPT-BASED XAI

Concept-based XAI answers the *what* question that attribution leaves open: instead of pointing to the inputs responsible for a prediction, it explains the prediction in terms of high-level, human-meaningful *concepts* (“striped”, “has wheels”, “expresses uncertainty”) of the kind a person would use to describe the same decision. Methods in this family differ chiefly in whether the concepts are built into the model

from the outset or recovered from a model that was trained without them.

2.2.2.1 *Explainable-by-design Models*

One way to guarantee that concepts genuinely take part in a model’s decision is to make them an explicit, intermediate stage of the computation, so that the prediction is forced to pass through a concept representation before reaching the output.

CONCEPT BOTTLENECK MODELS Concept Bottleneck Models (CBMs) [24] insert exactly such a stage: the network first maps the input to a set of human-specified concepts and then predicts the label from those concepts alone. Because the output depends on the input only through the concept layer, the model’s reasoning can be read off the activated concepts, and a user can intervene on them, correcting a mistaken concept and observing the effect on the prediction. The downside is supervision: training a CBM requires data annotated by human experts with concept labels, which is costly and ties the model to a fixed, predefined vocabulary.

CONCEPT EMBEDDING MODELS A known weakness of CBMs is an accuracy–interpretability trade-off: when the predefined concepts do not carry all the information the task requires, forcing the prediction through a narrow concept bottleneck degrades performance. Concept Embedding Models (CEMs) [13] loosen the bottleneck by representing each concept not as a single scalar but as a pair of high-dimensional embeddings, one for its active and one for its inactive state. This recovers accuracy comparable to an unconstrained black box while retaining concept-level interpretability and the possibility of intervention. Like CBMs, these models require a concept-annotated dataset.

LABEL-FREE CONCEPT BOTTLENECK MODELS Label-free CBMs (LF-CBMs) [30] address the annotation cost directly. Rather than relying on a hand-labelled concept set, they extract candidate concepts from an LLM and align the network’s activations with them through a vision–language model such as CLIP, converting a standard backbone into a CBM without any concept supervision. This makes the approach scalable to settings where dense concept annotation would be impractical.

CONCEPT BOTTLENECK LLMS The concept-bottleneck idea has recently been carried over to language. Concept Bottleneck LLMs (CB-LLMs) [44] insert a concept bottleneck layer into an LLM, routing text classification — and, in a generative variant, generation itself — through a set of interpretable concepts before the output. As in an LF-CBM, a hand-annotated concept set is not needed: the concepts

are generated by a prompted LLM and each training text is scored against them automatically, so the bottleneck is built without concept supervision. The explicit concept layer also makes the model *steerable*, since intervening on a concept predictably shifts its behaviour. Like the other methods in this group, CB-LLMs are *explainable by design*: interpretability is secured by training the model with the concept bottleneck built in.

Even so, that interpretability is only partial. In the generative variant the next-token prediction is shaped not by the concept bottleneck alone but by the bottleneck together with a parallel layer of unsupervised, uninterpretable neurons added to preserve generation quality. When the dataset labels are reused as concepts the bottleneck can be tiny — 4 neurons for the 4 AG News classes — yet in the authors’ implementation it runs alongside 768 unsupervised ones,¹ so only 4 of the 772 dimensions feeding each token prediction carry interpretable concepts.

2.2.2.2 *Post-hoc Methods*

The alternative to *explainable-by-design* models is to leave the trained model untouched and probe it for the concepts it has already come to represent.

TESTING WITH CONCEPT ACTIVATION VECTORS T-CAV [22] represents a concept as a direction in a layer’s activation space: given example inputs exhibiting the concept and a set of random ones, it fits a linear classifier to separate them and takes the normal to the decision boundary — the concept activation vector — as the concept’s direction. The model’s sensitivity to the concept is then quantified by the directional derivative of the output along this vector, yielding a global statement of how much a given concept influences a class prediction. The concept must still be supplied in advance, as a set of examples.

SAES AND MECHANISTIC INTERPRETABILITY A final, unsupervised strand removes even that requirement. Rather than testing for concepts named ahead of time, Sparse AutoEncoders (SAEs) learn a dictionary of concept-like features directly from a model’s internal activations [19], and have become a central tool of Mechanistic Interpretability (MechInt) — the effort to reverse-engineer the computations of neural networks, and of Large Language Models in particular. This unsupervised discovery of interpretable features is treated in full in section 2.3.3.

¹ The width of the unsupervised layer is fixed at 768 in the authors’ released code (<https://github.com/Trustworthy-ML-Lab/CB-LLMs>); the paper itself leaves it unspecified.

2.3 AUTOENCODERS

An autoencoder is a neural network trained to reconstruct its input through an intermediate latent representation [16, Ch 14]. It consists of two components: an **encoder** $f_\phi : \mathbb{R}^n \rightarrow \mathbb{R}^m$ that maps an input $\mathbf{x}$ to a latent code $\mathbf{z} = f_\phi(\mathbf{x})$, and a **decoder** $g_\theta : \mathbb{R}^m \rightarrow \mathbb{R}^n$ that reconstructs the input as $\hat{\mathbf{x}} = g_\theta(\mathbf{z})$. Training minimises a reconstruction loss, typically the mean squared error (MSE):

$$\mathcal{L}_{\text{recon}} = \text{MSE}(\mathbf{x}, \hat{\mathbf{x}}) = \|\mathbf{x} - \hat{\mathbf{x}}\|_2^2. \quad (2.2)$$

When $m < n$, the bottleneck forces the model to learn a compressed representation. When $m > n$, the network is overcomplete and additional constraints are required to prevent the trivial solution of learning the identity. In either case, the structure of the latent space is shaped by the choice of architecture and regularisation.

2.3.1 STANDARD AUTOENCODERS

2.3.1.1 Applications

IMAGE RETRIEVAL Krizhevsky and Hinton [26] proposed using deep autoencoders for content-based image retrieval. Their key insight is that a deep autoencoder can learn to map images to short binary codes that preserve semantic similarity: images that are perceptually similar are mapped to nearby codes in the latent space, allowing retrieval to be reduced to a fast lookup rather than an exhaustive search over the full database. They demonstrate this on a dataset of 1.6 million tiny images, showing that codes produced by deep autoencoders yield qualitatively and quantitatively better retrieval results than matching in raw pixel space. This work illustrated that the compressed latent representation learned by an autoencoder need not merely serve reconstruction, but can itself be a semantically meaningful and practically useful object.

ANOMALY DETECTION Sakurada and Yairi [40] were among the first to explicitly propose autoencoders as a tool for anomaly detection. Their core insight is that an autoencoder trained on normal data learns a compressed representation of the normal data distribution; at test time, inputs that deviate from this distribution are reconstructed poorly, and the resulting reconstruction error serves as an anomaly score. To substantiate this claim, they apply autoencoders to both synthetic data generated from the Lorenz system and real spacecraft telemetry, comparing against linear Principal Component Analysis (PCA) and kernel PCA. They demonstrate that autoencoders, by virtue of their nonlinear dimensionality reduction, are able to detect subtle anomalies that linear PCA misses, while remaining computationally lighter than kernel PCA. This work established reconstruction




Figure 2.2: Illustration of Denoising autoencoder model architecture, reproduced from Weng [52]. Note: this figure uses f for the decoder and g for the encoder, the reverse of the letter convention adopted in this thesis.

error as a natural unsupervised anomaly signal for autoencoder-based methods.

2.3.1.2 Denoising AutoEncoders

The denoising autoencoder was introduced by Vincent et al. [50] to avoid overfitting and improve the robustness of the learned representation. The input is partially corrupted by adding noise to or masking some values of the input vector in a stochastic manner, $\mathbf{x}_c \sim \mathcal{M}_{\mathcal{D}}(\mathbf{x}_c | \mathbf{x})$, as shown in fig. 2.2. The model is then trained to reconstruct the original uncorrupted input:

$$\mathcal{L}_{\text{DAE}} = \|\mathbf{x} - g_\theta(f_\phi(\mathbf{x}_c))\|_2^2. \quad (2.3)$$

This design is motivated by the fact that humans can easily recognise an object or a scene even when the view is partially occluded or corrupted. To “repair” the partially destroyed input, the denoising autoencoder has to discover and capture the relationship between dimensions of the input in order to infer missing pieces.

For high-dimensional input with high redundancy, like images, the model is likely to depend on evidence gathered from a combination of many input dimensions to recover the denoised version rather than to overfit one dimension. This builds up a good foundation for learning a *robust* latent representation.

The noise is controlled by a stochastic mapping $\mathcal{M}_{\mathcal{D}}(\mathbf{x}_c | \mathbf{x})$, and it is not specific to a particular type of corruption process (i.e., masking noise, Gaussian noise, salt-and-pepper noise, etc.). Naturally the corruption process can be equipped with prior knowledge [52].




Figure 2.3: Two-dimensional latent spaces learned on MNIST, coloured by digit class, reproduced from Shafkat [42]. Optimising for reconstruction alone (left) lets a standard autoencoder separate the classes into disconnected clusters with gaps that decode to unrealistic samples, whereas the Variational AutoEncoder objective (right) yields a smoother, more continuous arrangement, concentrated around the origin by the prior, that can be sampled and interpolated.

2.3.2 VARIATIONAL AUTOENCODERS

Variational AutoEncoders (VAEs) [23, 36] reinterpret the autoencoder as a probabilistic generative model. Rather than mapping an input to a single latent point, the encoder maps it to a *distribution* over latent codes, and the decoder is treated as a generative model that defines a distribution over inputs given a latent code. This probabilistic framing turns the autoencoder into a model from which new data can be sampled, and forces its latent space to have a smooth, continuous structure.

This smoothness is the property that sets a VAE apart from a standard autoencoder. Trained for reconstruction alone, an autoencoder is free to scatter its encodings into well-separated clusters with large gaps between them; those gaps are regions that decode to unrealistic outputs, so the space cannot be sampled or interpolated reliably. The VAE objective instead pushes encodings to fill the latent space more uniformly and to vary smoothly, so that nearby codes decode to similar inputs and paths between codes stay on the data manifold. Figure 2.3 illustrates the difference on a two-dimensional latent space trained on MNIST.

2.3.2.1 Variational Inference

Formally, a VAE models the data through a generative process in which a latent code is first drawn from a prior $\mathbf{z} \sim p(\mathbf{z})$ — typically

a standard Gaussian $p(\mathbf{z}) = \mathcal{N}(\mathbf{0}, \mathbf{I})$ — and an observation is then drawn from a decoder distribution $p_\theta(\mathbf{x} | \mathbf{z})$ parameterised by a neural network. Training would ideally maximise the marginal likelihood

$$p_\theta(\mathbf{x}) = \int p_\theta(\mathbf{x} | \mathbf{z}) p(\mathbf{z}) d\mathbf{z},$$

but this integral, and the corresponding posterior $p_\theta(\mathbf{z} | \mathbf{x}) = p_\theta(\mathbf{x} | \mathbf{z})p_\theta(\mathbf{z})/p_\theta(\mathbf{x})$, are intractable for any non-trivial decoder.

The key idea of Kingma and Welling [23] is to sidestep this intractability with *amortised variational inference*: an encoder network $q_\phi(\mathbf{z} | \mathbf{x})$ is trained to approximate the true posterior, with all data points sharing the same inference parameters ϕ . The encoder outputs the parameters of a diagonal Gaussian,

$$q_\phi(\mathbf{z} | \mathbf{x}) = \mathcal{N}(\mathbf{z}; \boldsymbol{\mu}_\phi(\mathbf{x}), \text{diag}(\sigma_\phi^2(\mathbf{x}))), \quad (2.4)$$

so that each input is mapped to a mean and a (diagonal) covariance rather than to a single point.

Instead of the intractable likelihood, the VAE maximises a tractable lower bound on it, the Evidence Lower Bound (ELBO):

$$\log p_\theta(\mathbf{x}) \geq \underbrace{\mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x} | \mathbf{z})]}_{\text{reconstruction}} - \underbrace{D_{\text{KL}}(q_\phi(\mathbf{z} | \mathbf{x}) \| p(\mathbf{z}))}_{\text{regularisation}}. \quad (2.5)$$

The bound follows from the identity

$$\log p_\theta(\mathbf{x}) = \text{ELBO} + D_{\text{KL}}(q_\phi(\mathbf{z} | \mathbf{x}) \| p_\theta(\mathbf{z} | \mathbf{x})),$$

together with the non-negativity of the Kullback–Leibler (KL) divergence; it is tight precisely when the approximate posterior matches the true posterior $p_\theta(\mathbf{z} | \mathbf{x})$. Maximising the ELBO simultaneously pushes the decoder to assign high likelihood to the data (the reconstruction term) and keeps the approximate posterior close to the prior (the KL divergence term). Negating it yields the training loss:

$$\mathcal{L}_{\text{VAE}} = \mathcal{L}_{\text{recon}} + D_{\text{KL}}(q_\phi(\mathbf{z} | \mathbf{x}) \| p(\mathbf{z})). \quad (2.6)$$

When the decoder is Gaussian with fixed isotropic variance, the first term reduces, up to an additive constant, to the MSE reconstruction loss of eq. (2.2); the second term, being the divergence between two Gaussians, has a closed form. For the standard-normal prior it becomes:

$$D_{\text{KL}}(q_\phi(\mathbf{z} | \mathbf{x}) \| p(\mathbf{z})) = \frac{1}{2} \sum_{j=1}^m (\mu_j^2 + \sigma_j^2 - \log \sigma_j^2 - 1),$$

where μ_j and σ_j are the components of the encoder's predicted mean and standard deviation.

Evaluating the reconstruction term requires sampling $\mathbf{z}$ from $q_\phi(\mathbf{z} | \mathbf{x})$, an operation that is not differentiable with respect to

the encoder parameters and would therefore block gradient-based training. The VAE resolves this with the *reparameterisation trick*: the stochasticity is moved to an auxiliary noise variable, and the sample is expressed as a deterministic, differentiable function of the encoder outputs,

$$\mathbf{z} = \boldsymbol{\mu}_\phi(\mathbf{x}) + \boldsymbol{\sigma}_\phi(\mathbf{x}) \odot \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}),$$

where $\odot$ denotes the element-wise product. Gradients can now flow through $\boldsymbol{\mu}_\phi$ and $\boldsymbol{\sigma}_\phi$ while the randomness is confined to $\boldsymbol{\epsilon}$, making the whole objective trainable end-to-end by stochastic gradient descent.

2.3.2.2 Architectures

There exist several variants that modify the base VAE to shape the latent space or the form of the latent code; two are particularly relevant here.

β -VAE Higgins et al. [18] introduced a single modification: a scalar weight β on the KL term of the objective,

$$\mathcal{L}_{\beta\text{-VAE}} = \mathcal{L}_{\text{recon}} + \beta D_{\text{KL}}(q_\phi(\mathbf{z} | \mathbf{x}) \| p(\mathbf{z})). \quad (2.7)$$

Setting $\beta > 1$ places additional pressure on the approximate posterior to match the factorised prior, which encourages individual latent dimensions to capture statistically independent factors of variation — a property known as *disentanglement*. The standard VAE is recovered at $\beta = 1$. Changing β offers a trade-off: a larger value improves the structure and interpretability of the latent space at the cost of reconstruction fidelity.

VQ-VAE van den Oord, Vinyals and Kavukcuoglu [48] replaced the continuous Gaussian latent with a *discrete* one. The encoder output is quantised to the nearest entry of a learned codebook $\{\mathbf{e}_1, \dots, \mathbf{e}_K\}$, so that each latent is represented by a discrete index rather than a continuous vector. This sidesteps the *posterior collapse* that can affect continuous VAEs paired with expressive autoregressive decoders, in which the decoder learns to ignore the latent entirely. Because the nearest-neighbour lookup is non-differentiable, gradients are passed to the encoder with a straight-through estimator and the codebook is learned through dedicated loss terms. Vector-Quantized Variational AutoEncoder (VQ-VAE) underpins many modern generative systems, providing the discrete latent vocabulary over which an autoregressive prior is subsequently trained.

2.3.2.3 Applications

VAEs are used both as practical components in large generative systems and for interpretable representation learning; the latter is the more relevant to this thesis.

GENERATIVE MODELLING Although a VAE can act as a standalone generator, in modern systems it more often serves as a *compression backbone* for other generative models: latent diffusion models run the diffusion process in the compact latent space of an autoencoder rather than in pixel space. The autoencoder behind Stable Diffusion [38] is regularised only weakly, with a small KL penalty (in the manner of a VAE) or with vector quantisation. Discrete VQ-VAEs [48] encode data into a grid of discrete codes from a learned codebook, which an autoregressive prior can then generate. The smooth latent also serves as a search space: optimising over it has been used to search for new molecules with desired properties [15].

ANOMALY DETECTION Like standard autoencoders, VAEs can flag anomalies, but they replace the single reconstruction error with a probabilistic score: the reconstruction probability, or the ELBO, gives a principled measure of how well the model explains an input [3]. This is the probabilistic counterpart of the approach described in section 2.3.1.1.

REPRESENTATION LEARNING VAEs are also a standard tool for unsupervised representation learning. The β -VAE (section 2.3.2.2) is the canonical example: up-weighting the KL term pressures the model towards *disentangled* latents whose individual dimensions capture independent, semantically meaningful factors of variation [18], yielding structured low-dimensional features for downstream use.

INTERPRETABILITY The structured latent that makes VAEs good representation learners also makes them attractive for explainability. Their smoothness supports *counterfactual explanation* (section 2.2.1.2), as perturbing the latent yields plausible, on-manifold variants of an input [20]. A substantial line of work uses VAEs for *causal representation learning*, aligning latent factors with underlying causal variables [17, 25]. VAEs have also been paired with concept-based models such as self-explaining networks [1]. Closest to this thesis, variational ideas have been brought to sparse dictionary learning — variational sparse coding [47] and variational Sparse AutoEncoders [4] — seeking to combine the interpretability of sparse codes with the structured latent of a VAE.

2.3.3 SPARSE AUTOENCODERS

Sparse AutoEncoders (SAEs) are one of the main variants of autoencoders. Just like autoencoders, they learn a dictionary of feature directions from the input data, but they typically are overcomplete ($m \gg n$)




Figure 2.4: Neuron 4e:55, a polysemantic neuron from InceptionV1 which responds to cat faces, fronts of cars, and cat legs. Reproduced from Olah et al. [31].

and their latent representation is forced to be sparse using some kind of sparsity mechanism. Given input $\mathbf{x} \in \mathbb{R}^n$, the SAE encodes it as:

$$\mathbf{z} = \text{ReLU}(\mathbf{W}_e(\mathbf{x} - \mathbf{b}_{\text{pre}}) + \mathbf{b}_e), \quad (2.8)$$

where $\mathbf{W}_e \in \mathbb{R}^{m \times n}$ is the encoder weight matrix, $\mathbf{b}_e \in \mathbb{R}^m$ is the encoder bias, and $\mathbf{b}_{\text{pre}} \in \mathbb{R}^n$ is a pre-encoder bias subtracted from the input before encoding. This learned offset allows the encoder to operate on roughly mean-centred activations, improving training stability [6]. The decoder then reconstructs the input as:

$$\hat{\mathbf{x}} = \mathbf{W}_d \mathbf{z} + \mathbf{b}_{\text{pre}}, \quad (2.9)$$

where $\mathbf{W}_d \in \mathbb{R}^{n \times m}$ is the decoder weight matrix. Each column $\mathbf{d}_i$ of $\mathbf{W}_d$ is a dictionary atom, representing a candidate feature direction in activation space.

When used in the context of MechInt, the overcompleteness and the sparsity constraint are motivated by the *superposition hypothesis*.

2.3.3.1 The Superposition Hypothesis

A central challenge in MechInt is that individual neurons in neural networks tend to be *polysemantic*: they activate in response to multiple, seemingly unrelated concepts [12, 31]. An example of this can be seen in fig. 2.4. This runs contrary to the intuition that an interpretable model should have neurons with clean, well-defined roles.

Elhage et al. [12] proposed the *superposition hypothesis* to explain this phenomenon. The core idea is that a network may represent more features than it has dimensions by exploiting the fact that natural inputs are sparse: at any given time, only a small fraction of all possible features are relevant. Under this condition, a network can pack many near-orthogonal feature directions into a lower-dimensional space with limited interference; polysemanticity ceases to be a failure of the model, and instead becomes a rational compression strategy under capacity constraints.

This has a direct implication for interpretability: the natural basis of a model’s activations is not the neuron basis, but a larger, overcomplete set of feature directions embedded within it. Recovering this basis requires moving beyond individual neuron analysis, towards methods that can identify sparse, distributed structure. This is precisely the goal of *sparse dictionary learning* [32], a classical technique from computational neuroscience that SAEs adapt to the setting of modern language models.

2.3.3.2 Architectures

There are many ways to enforce sparsity, each with its pros and cons. Here is a selection of methods that range from the initial SAE idea to more modern architectures that try to improve on it.

VANILLA ℓ_1 This is the original architecture proposed by Huben et al. [19]. It uses an ℓ_1 penalty² on the latent code to encourage most of $\mathbf{z}$ ’s entries to be exactly zero, so that any given activation is reconstructed from a small number of active features. The loss used to train this model is the following:

$$\mathcal{L} = \mathcal{L}_{\text{recon}} + \lambda \|\mathbf{z}\|_1. \quad (2.10)$$

The scalar λ controls the sparsity–reconstruction trade-off and must be tuned carefully: too large, and reconstruction degrades; too small, and the learned features lose monosemanticity. A known side effect of ℓ_1 regularisation is *shrinkage* [46, 53]: beyond driving inactive features to zero, the penalty also biases active feature magnitudes below their reconstruction-optimal values, causing the model to systematically underestimate the true activation of each feature [34, 53]. Constraining the decoder columns to unit ℓ_2 norm after each gradient step partially mitigates this, by preventing the model from reducing the ℓ_1 penalty through the trivial reparameterisation of scaling down encoder outputs and compensating with larger decoder norms [34]. There is, however, a residual gradient-level bias that is structurally entangled with the ℓ_1 penalty and cannot be resolved with this trick.

TOPK SAEs Gao et al. [14] proposed replacing the ℓ_1 penalty with a hard sparsity constraint, implemented via a TopK activation function that retains only the k largest pre-activations and zeros the rest. This directly controls the number of active features per forward pass, eliminating the need to tune λ and removing the shrinkage bias entirely. The full training objective is:

$$\mathcal{L} = \mathcal{L}_{\text{recon}} + \alpha \mathcal{L}_{\text{aux}}, \quad (2.11)$$

² The ℓ_p norm is defined as $\ell_p(\mathbf{x}) = \|\mathbf{x}\|_p = (\sum_i |x_i|^p)^{1/p}$, $p \geq 1$.

where the auxiliary loss $\mathcal{L}_{\text{aux}}$ addresses the problem of dead latents, i. e., features that cease to activate entirely during training. Latents are flagged as dead if they have not activated for a predetermined number of tokens. Given the reconstruction residual $\mathbf{e} = \mathbf{x} - \hat{\mathbf{x}}$, the top- k_{aux} dead latents are selected and used to reconstruct $\mathbf{e}$; the resulting reconstruction error is $\mathcal{L}_{\text{aux}}$. This gives dead latents a gradient signal proportional to how much of the residual they could explain, encouraging them to activate rather than remain permanently dormant. Gao et al. [14] demonstrate that TopK SAEs achieve better reconstruction at equivalent sparsity levels compared to ℓ_1 baselines, making them a strong and widely used alternative.

GATED SAES Rajamanoharan et al. [34] introduced an architecture that separates the decision of whether a feature is active from the decision of how strongly it activates. A dedicated gating pathway produces a binary mask via a Heaviside step function (approximated during training with a straight-through estimator) which is applied element-wise to the encoder output. This architectural separation allows the model to learn feature presence and magnitude more independently, addressing shrinkage from a structural rather than a loss-based perspective.

JUMPRELU SAES Rajamanoharan et al. [35] proposed replacing the standard ReLU with a JumpReLU activation, which is zero below a learned per-feature threshold θ_i and equal to the pre-activation above it:

$$\text{JumpReLU}_{\theta_i}(z_i) = z_i \cdot \mathbb{1}[z_i > \theta_i]. \quad (2.12)$$

The key advantage of this parameterisation is that it provides a dedicated, per-feature learnable threshold, which allows the ℓ_0 norm³ of the feature activations — $\sum_i H(\pi_i - \theta_i)$, where π_i are the encoder pre-activations and $H(x)$ is the Heaviside step function — to be trained directly via a straight-through estimator applied to the threshold parameter θ_i . Because each threshold can be positioned near the actual boundary between active and inactive pre-activations for its feature, the gradient signal is stable and controllable, making ℓ_0 training practical in a way that would not be possible with a fixed threshold. Training with ℓ_0 rather than ℓ_1 eliminates shrinkage entirely, since ℓ_0 penalises only the number of active features and not their intensity. JumpReLU SAEs achieve comparable performance to TopK SAEs.

2.3.3.3 Applications

Unlike general-purpose autoencoders and VAEs, which are applied across a range of domains, SAEs find their primary application in

³ The ℓ_0 norm is a pseudo-norm defined as the number of non-zero vector components: $\ell_0(\mathbf{x}) = \sum_i \mathbb{1}[x_i > 0]$.

interpreting the internal activations of neural networks, and [MechInt](#) remains their most prominent application, making it the focus of this section. [SAEs](#) can be trained on activations from any layer or component of a transformer (section 2.1.2); common targets include the residual stream at various depths, [MLP](#) layer outputs, and attention head outputs [6, 19]. The resulting dictionary then supports the applications described below.

FEATURE DISCOVERY The primary use of [SAEs](#) is *feature discovery*: training [an SAE](#) on a model’s internal activations and examining the resulting dictionary atoms to identify human-interpretable concepts. The underlying hypothesis is that each atom, active for only a semantically coherent set of inputs, corresponds to a monosemantic feature of the model’s computation — a clean unit of meaning that the model had obscured in the neuron basis. In practice these features often correspond to recognisable concepts: Bricken et al. [6] trained [an SAE](#) on the [MLP](#) of a one-layer transformer and found features responsive to specific tokens, syntactic roles, and semantic categories, while at production scale Templeton et al. [45] report features for famous people, countries, emotions, or even biases (like racism, sexism, hatred) and sarcastic praise. These results provide strong empirical support for the superposition hypothesis and suggest that [SAEs](#) can serve as a practical lens through which to study what information [LLMs](#) represent.

STEERING The features recovered by [an SAE](#) are not merely descriptive: because each corresponds to a direction in activation space, its activation can be clamped to an arbitrary value during a forward pass to causally *steer* the model’s behaviour. Templeton et al. [45] demonstrated this by amplifying a single feature, causing the model to compulsively reference the Golden Gate Bridge regardless of the input. Such interventions show that the discovered features are causally implicated in the model’s computation rather than merely correlated with its inputs, and open a path to controllable generation and targeted safety interventions.

CIRCUIT ANALYSIS Beyond cataloguing individual features, [SAEs](#) can serve as the building blocks for reconstructing the *circuits* a model uses to perform a computation, expressing the interactions between features across layers as an interpretable causal graph [28]. More recent work favours *transcoders* — a variant trained to approximate the input–output behaviour of an [MLP](#) sublayer with a sparse hidden representation, rather than to reconstruct a single activation — which should disentangle the model’s computation more cleanly and yield sparser, more faithful circuits [2, 10].

2.3.3.4 Evaluation

Assessing the quality of learned features is non-trivial, as interpretability is inherently subjective. Several proxy metrics are used in practice. *Reconstruction quality* is judged by its effect on the model rather than by the raw reconstruction error: the SAE’s output is substituted for the original activations, the model is run forward, and the resulting increase in its next-token CE loss is measured [14]. A faithful SAE leaves this loss almost unchanged, indicating that it has retained the information the model actually uses. *Sparsity* is quantified via the average ℓ_0 norm of the latent codes, i.e., the mean number of active features per token. *Automated interpretability* pipelines use a second language model to generate and score natural-language descriptions of individual features [5], providing a scalable if imperfect proxy for human judgement.

CONCLUSIONS

This chapter has assembled the three strands the thesis builds on: the transformer and its residual stream (section 2.1), whose dense, polysemantic activations are the object to be decomposed; the shift in explainability from attribution to concepts (section 2.2), and the concept-based methods that describe a model’s computation in human-meaningful terms; and the autoencoder family (section 2.3), from the variational and discrete latents of VAEs and VQ-VAEs to the overcomplete, sparse dictionaries of SAEs.

A single limitation runs through these concept-based methods. Whether a concept is a unit in the bottleneck of a CBM or an atom in the dictionary of an SAE, it is encoded as a single scalar and can vary only in magnitude. This is what forces the sparsity–reconstruction trade-off at the heart of the SAE literature — a faithful reconstruction needs many active concepts, an interpretable decomposition only few — and the architectures of section 2.3.3.2 all work within it. CEMs are the exception already met: by giving each concept a vector embedding rather than a scalar, they recover the capacity a scalar bottleneck gives up. Section 3.3 carries that idea into the unsupervised setting of activation decomposition — combining the vector-valued concepts of CEMs with the variational inference and discrete latents developed in section 2.3 — and derives the resulting model from first principles.

METHODS

3.1 AIM OF THE WORK

SAEs have become the standard tool for decomposing the activations of an LLM into a dictionary of concept-like features (section 2.3.3). Their reach is bounded, though, by a structural trade-off between sparsity and reconstruction [14]: because each concept enters the reconstruction through a single scalar coefficient, reconstructing an activation faithfully tends to require many concepts active at once, whereas the interpretability that motivates the decomposition rests on keeping only a few active per token [6, 19]. Pushed to a useful sparsity, the reconstruction degrades far enough that the reconstructed activation leaves the data manifold. This is more than an aesthetic flaw: a downstream circuit then processes an out-of-distribution signal even before any intervention, and when a feature is amplified or ablated the resulting change in behaviour conflates the intended edit with this uncontrolled reconstruction error, so the very analyses the decomposition is built for become less reliable than they appear.

The limitation lies in the representation rather than the training: a scalar can record only how strongly a concept is present. This thesis asks whether giving each concept a *vector* embedding relaxes the trade-off. The idea is familiar from CEMs (section 2.2.2), which replace the scalar concept bottleneck of a CBM with per-concept embeddings and so recover the accuracy a narrow bottleneck gives up [13] — but CEMs are supervised, trained against concept labels. The aim here is to carry the vector-concept idea into the *unsupervised* setting of activation decomposition, and to do so from first principles, defining a probabilistic generative model whose inference procedure and training objective follow from the model itself rather than being assembled by hand. The starting point is the Learnable Concept-based Model (LCBM) of De Santis et al. [8], an unsupervised concept-based image classifier whose encoder produces a vector embedding per concept, gates each concept by comparing that embedding to a learnable prototype, and reconstructs the input and simultaneously classifies it from the embeddings of the active concepts. LCBM supplied the idea rather than the architecture: this thesis keeps its core structure — concepts as embeddings gated against prototypes and decoded by reconstruction — but rebuilds it from the ground up as a fully variational generative model, giving the concept embeddings a probabilistic latent that LCBM leaves deterministic and dropping the supervised classification task, so that the model reconstructs activations from no labels at all.

sure about this?

That model — the Variational Auto Embedding Encoder (VAEE) — sits where three of the families surveyed in chapter 2 meet. From CEMs it takes the representation of a concept as a high-dimensional embedding rather than a scalar; from VAEs it takes a probabilistic latent and amortised variational inference (section 2.3.2); and from VQ-VAEs it takes a learned set of concept prototypes together with a discrete choice of which concepts are active (section 2.3.2.2). The result is a probabilistic counterpart to discrete quantisation: independent Bernoulli activations route continuous, prototype-anchored embeddings into separate concept slots, retaining the expressivity of a continuous latent while keeping the routing sparse and selective. Nothing in this construction is specific to language models — the vector-concept primitive applies wherever a scalar concept code trades capacity for simplicity — but this thesis develops and demonstrates it for one application: the decomposition of LLM activations.

What this thesis contributes, then, is the architecture itself and its derivation: an unsupervised, vector-valued concept model obtained from a probabilistic generative model rather than assembled by hand. Its aim is correspondingly bounded — to expand the representational capacity of a concept code, leaving the interpretability and steering of the recovered concepts to future work — and chapter 4 tests it on the decomposition of LLM activations.

3.2 DEFINITIONS / NOTATION

3.3 MODEL DEVELOPMENT

3.3.1 PROBABILISTIC FORMULATION

3.3.2 ARCHITECTURES

CONCLUSIONS

EXPERIMENTS

4.1 MODEL TRAINING

4.2 EVALUATION METRICS

4.2.1 MSE

4.2.2 ℓ_0

4.2.3 PERPLEXITY

4.3 RESULTS

CONCLUSIONS

CONCLUSIONS

5.1 CONTRIBUTIONS

5.2 FUTURE WORK



VAEE DERIVATIONS

This appendix collects the full derivations supporting the probabilistic formulation of section 3.3.1: the optimal form of the gating mechanism (appendix A.1) and the explicit derivation of the evidence lower bound (appendix A.2).

A.1 OPTIMAL GATING MECHANISM

A.2 EXPLICIT ELBO DERIVATION

BIBLIOGRAPHY

- [1] David Alvarez Melis and Tommi Jaakkola. ‘Towards Robust Interpretability with Self-Explaining Neural Networks’. In: *Advances in Neural Information Processing Systems*. Vol. 31. Curran Associates, Inc., 2018. URL: https://papers.nips.cc/paper_files/paper/2018/hash/3e9f0fc9b2f89e043bc6233994dfcf76-Abstract.html (visited on 11/06/2026).
- [2] Emmanuel Ameisen et al. *Circuit Tracing: Revealing Computational Graphs in Language Models*. Tech. rep. Transformer Circuits Thread, Mar. 2025. URL: <https://transformer-circuits.pub/2025/attribution-graphs/methods.html> (visited on 11/06/2026).
- [3] Jinwon An and Sungzoon Cho. *Variational Autoencoder Based Anomaly Detection Using Reconstruction Probability*. Dec. 2015. URL: <http://dm.snu.ac.kr/static/docs/TR/SNUDM-TR-2015-03.pdf>.
- [4] Zachary Baker and Yuxiao Li. *Analysis of Variational Sparse Autoencoders*. Comment: 15 pages, 11 figures. Oct. 2025. DOI: [10.48550/arXiv.2509.22994](https://doi.org/10.48550/arXiv.2509.22994). arXiv: [2509.22994](https://arxiv.org/abs/2509.22994) [cs]. URL: <http://arxiv.org/abs/2509.22994> (visited on 20/10/2025).
- [5] Steven Bills, Nick Cammarata, Dan Mossing, Henk Tillman, Leo Gao, Gabriel Goh, Ilya Sutskever, Jan Leike, Jeff Wu and William Saunders. *Language Models Can Explain Neurons in Language Models*. Tech. rep. OpenAI, May 2023. URL: <https://openaipublic.blob.core.windows.net/neuron-explainer/paper/index.html> (visited on 11/06/2026).
- [6] Trenton Bricken et al. *Towards Monosemanticity: Decomposing Language Models With Dictionary Learning*. Tech. rep. Transformer Circuits Thread, Oct. 2023. URL: <https://transformer-circuits.pub/2023/monosemantic-features/index.html> (visited on 11/06/2026).
- [7] Tom Brown et al. ‘Language Models Are Few-Shot Learners’. In: *Advances in Neural Information Processing Systems*. Vol. 33. Curran Associates, Inc., 2020, pp. 1877–1901. URL: <https://proceedings.neurips.cc/paper/2020/hash/1457c0d6bfc4967418bfb8ac142f64a-Abstract.html> (visited on 13/06/2026).
- [8] Francesco De Santis, Philippe Bich, Gabriele Ciravegna, Pietro Barbiero, Tania Cerquitelli and Danilo Giordano. ‘Towards Better Generalization and Interpretability in Unsupervised Concept-

- Based Models’. In: *Machine Learning and Knowledge Discovery in Databases. Research Track: European Conference, ECML PKDD 2025, Porto, Portugal, September 15–19, 2025, Proceedings, Part III*. Berlin, Heidelberg: Springer-Verlag, Oct. 2025, pp. 478–494. ISBN: 978-3-032-06065-5. DOI: [10.1007/978-3-032-06066-2_28](https://doi.org/10.1007/978-3-032-06066-2_28). URL: https://doi.org/10.1007/978-3-032-06066-2_28 (visited on 11/06/2026).
- [9] Qingxiu Dong et al. ‘A Survey on In-context Learning’. In: *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Ed. by Yaser Al-Onaizan, Mohit Bansal and Yun-Nung Chen. Miami, Florida, USA: Association for Computational Linguistics, Nov. 2024, pp. 1107–1128. DOI: [10.18653/v1/2024.emnlp-main.64](https://doi.org/10.18653/v1/2024.emnlp-main.64). URL: <https://aclanthology.org/2024.emnlp-main.64/> (visited on 11/06/2026).
- [10] Jacob Dunefsky, Philippe Chlenski and Neel Nanda. ‘Transcoders Find Interpretable LLM Feature Circuits’. In: *Advances in Neural Information Processing Systems* 37 (Dec. 2024), pp. 24375–24410. DOI: [10.52202/079017-0768](https://doi.org/10.52202/079017-0768). URL: https://proceedings.neurips.cc/paper_files/paper/2024/hash/2b8f4db0464cc5b6e9d5e6bea4b9f308-Abstract-Conference.html (visited on 10/06/2026).
- [11] Nelson Elhage et al. *A Mathematical Framework for Transformer Circuits*. Tech. rep. Transformer Circuits Thread, Dec. 2021. URL: <https://transformer-circuits.pub/2021/framework/index.html> (visited on 13/06/2026).
- [12] Nelson Elhage et al. *Toy Models of Superposition*. Tech. rep. Transformer Circuits Thread, Sept. 2022. URL: https://transformer-circuits.pub/2022/toy_model/index.html (visited on 11/06/2026).
- [13] Mateo Espinosa Zarlenga et al. ‘Concept Embedding Models: Beyond the Accuracy-Explainability Trade-Off’. In: *Advances in Neural Information Processing Systems* 35 (Dec. 2022), pp. 21400–21413. URL: https://proceedings.neurips.cc/paper_files/paper/2022/hash/867c06823281e506e8059f5c13a57f75-Abstract-Conference.html (visited on 12/06/2026).
- [14] Leo Gao, Tom Dupre la Tour, Henk Tillman, Gabriel Goh, Rajan Troll, Alec Radford, Ilya Sutskever, Jan Leike and Jeffrey Wu. ‘Scaling and Evaluating Sparse Autoencoders’. In: *The Thirteenth International Conference on Learning Representations*. Oct. 2024. URL: <https://openreview.net/forum?id=tcsZt9ZNKD> (visited on 14/05/2026).
- [15] Rafael Gómez-Bombarelli, Jennifer N. Wei, David Duvenaud, José Miguel Hernández-Lobato, Benjamín Sánchez-Lengeling, Dennis Sheberla, Jorge Aguilera-Iparraguirre, Timothy D. Hirzel,

- Ryan P. Adams and Alán Aspuru-Guzik. 'Automatic Chemical Design Using a Data-Driven Continuous Representation of Molecules'. In: *ACS Central Science* 4.2 (Feb. 2018), pp. 268–276. ISSN: 2374-7943. DOI: [10.1021/acscentsci.7b00572](https://doi.org/10.1021/acscentsci.7b00572). URL: <https://doi.org/10.1021/acscentsci.7b00572> (visited on 10/06/2026).
- [16] Ian Goodfellow, Aaron Courville and Yoshua Bengio. *Deep Learning*. Adaptive Computation and Machine Learning. "Deep learning is a form of machine learning that enables computers to learn from experience and understand the world in terms of a hierarchy of concepts. Because the computer gathers knowledge from experience, there is no need for a human computer operator to formally specify all the knowledge that the computer needs. The hierarchy of concepts allows the computer to learn complicated concepts by building them out of simpler ones; a graph of these hierarchies would be many layers deep. This book introduces a broad range of topics in deep learning. The text offers mathematical and conceptual background, covering relevant concepts in linear algebra, probability theory and information theory, numerical computation, and machine learning. It describes deep learning techniques used by practitioners in industry, including deep feedforward networks, regularization, optimization algorithms, convolutional networks, sequence modeling, and practical methodology; and it surveys such applications as natural language processing, speech recognition, computer vision, online recommendation systems, bioinformatics, and video games. Finally, the book offers research perspectives, covering such theoretical topics as linear factor models, autoencoders, representation learning, structured probabilistic models, Monte Carlo methods, the partition function, approximate inference, and deep generative models. Deep Learning can be used by undergraduate or graduate students planning careers in either industry or research, and by software engineers who want to begin using deep learning in their products or platforms. A website offers supplementary material for both readers and instructors"—Publisher's description Introduction – Applied math and machine learning basics. Linear algebra – Probability and information theory – Numerical computation – Machine learning basics – Deep networks: modern practices. Deep feedforward networks – Regularization for deep learning – Optimization for training deep models – Convolutional networks – Sequence modeling: recurrent and recursive nets – Practical methodology – Applications – Deep learning research. Linear factor models – Autoencoders – Representation learning – Structured probabilistic models for deep learning – Monte Carlo

- methods – Confronting the partition function – Approximate inference – Deep generative models
Includes bibliographical references and index. Cambridge, Massachusetts: The MIT Press, 2016. ISBN: 978-0-262-03561-3 978-0-262-33737-3.
- [17] Yash Goyal, Amir Feder, Uri Shalit and Been Kim. *Explaining Classifiers with Causal Concept Effect (CaCE)*. Feb. 2020. DOI: [10.48550/arXiv.1907.07165](https://doi.org/10.48550/arXiv.1907.07165). arXiv: [1907.07165](https://arxiv.org/abs/1907.07165) [cs.LG]. URL: <http://arxiv.org/abs/1907.07165> (visited on 01/06/2026).
 - [18] Irina Higgins, Loic Matthey, Arka Pal, Christopher Burgess, Xavier Glorot, Matthew Botvinick, Shakir Mohamed and Alexander Lerchner. ‘Beta-VAE: Learning Basic Visual Concepts with a Constrained Variational Framework’. In: *International Conference on Learning Representations*. Feb. 2017. URL: <https://openreview.net/forum?id=Sy2fzU9gl> (visited on 10/06/2026).
 - [19] Robert Huben, Hoagy Cunningham, Logan Riggs Smith, Aidan Ewart and Lee Sharkey. ‘Sparse Autoencoders Find Highly Interpretable Features in Language Models’. In: *The Twelfth International Conference on Learning Representations*. Oct. 2023. URL: <https://openreview.net/forum?id=F76bwRSLeK> (visited on 28/05/2026).
 - [20] Shalmali Joshi, Oluwasanmi Koyejo, Warut Vijitbenjaronk, Been Kim and Joydeep Ghosh. *Towards Realistic Individual Recourse and Actionable Explanations in Black-Box Decision Making Systems*. July 2019. DOI: [10.48550/arXiv.1907.09615](https://doi.org/10.48550/arXiv.1907.09615). arXiv: [1907.09615](https://arxiv.org/abs/1907.09615) [cs.LG]. URL: <http://arxiv.org/abs/1907.09615> (visited on 11/06/2026).
 - [21] Been Kim, Rajiv Khanna and Oluwasanmi O Koyejo. ‘Examples Are Not Enough, Learn to Criticize! Criticism for Interpretability’. In: *Advances in Neural Information Processing Systems*. Vol. 29. Curran Associates, Inc., 2016. URL: https://proceedings.neurips.cc/paper_files/paper/2016/hash/5680522b8e2bb01943234bce7bf84534-Abstract.html (visited on 11/06/2026).
 - [22] Been Kim, Martin Wattenberg, Justin Gilmer, Carrie Cai, James Wexler, Fernanda Viegas and Rory Sayres. ‘Interpretability Beyond Feature Attribution: Quantitative Testing with Concept Activation Vectors (TCAV)’. In: *Proceedings of the 35th International Conference on Machine Learning*. PMLR, July 2018, pp. 2668–2677. URL: <https://proceedings.mlr.press/v80/kim18d.html> (visited on 11/06/2026).
 - [23] Diederik P. Kingma and Max Welling. ‘Auto-Encoding Variational Bayes’. In: *Useful for the reparametrization trick (sec 2.4)*. International Conference on Learning Representations, May 2014. DOI: [10.48550/arXiv.1312.6114](https://doi.org/10.48550/arXiv.1312.6114). arXiv: [1312.6114](https://arxiv.org/abs/1312.6114)

- [stat]. URL: <http://arxiv.org/abs/1312.6114> (visited on 07/11/2025).
- [24] Pang Wei Koh, Thao Nguyen, Yew Siang Tang, Stephen Mussmann, Emma Pierson, Been Kim and Percy Liang. ‘Concept Bottleneck Models’. In: *Proceedings of the 37th International Conference on Machine Learning*. PMLR, Nov. 2020, pp. 5338–5348. URL: <https://proceedings.mlr.press/v119/koh20a.html> (visited on 11/06/2026).
- [25] Aneesh Komanduri, Yongkai Wu, Wen Huang, Feng Chen and Xintao Wu. ‘SCM-VAE: Learning Identifiable Causal Representations via Structural Knowledge’. In: *2022 IEEE International Conference on Big Data (Big Data)*. Dec. 2022, pp. 1014–1023. DOI: [10.1109/BigData55660.2022.10021114](https://doi.org/10.1109/BigData55660.2022.10021114). URL: <https://ieeexplore.ieee.org/document/10021114> (visited on 21/04/2026).
- [26] Alex Krizhevsky and Geoffrey E. Hinton. ‘Using Very Deep Autoencoders for Content-Based Image Retrieval’. In: *19th European Symposium on Artificial Neural Networks, Computational Intelligence and Machine Learning* (Apr. 2011).
- [27] Scott M Lundberg and Su-In Lee. ‘A Unified Approach to Interpreting Model Predictions’. In: *Advances in Neural Information Processing Systems*. Vol. 30. Curran Associates, Inc., 2017. URL: https://proceedings.neurips.cc/paper_files/paper/2017/hash/8a20a8621978632d76c43dfd28b67767-Abstract.html (visited on 01/06/2026).
- [28] Samuel Marks, Can Rager, Eric Michaud, Yonatan Belinkov, David Bau and Aaron Mueller. ‘Sparse Feature Circuits: Discovering and Editing Interpretable Causal Graphs in Language Models’. In: *International Conference on Learning Representations 2025* (May 2025), pp. 23888–23923. URL: https://proceedings.iclr.cc/paper_files/paper/2025/hash/3ba4d47a83e498c2b1a0868cba20f6de-Abstract-Conference.html (visited on 10/06/2026).
- [29] Christoph Molnar. *Interpretable Machine Learning: A Guide for Making Black Box Models Explainable*. 3rd edition. Munich, Germany: Christoph Molnar, 2025. ISBN: 978-3-911578-03-5.
- [30] Tuomas Oikarinen, Subhro Das, Lam M. Nguyen and Tsui-Wei Weng. ‘Label-Free Concept Bottleneck Models’. In: *The Eleventh International Conference on Learning Representations*. Sept. 2022. URL: <https://openreview.net/forum?id=FLCg47MNvBA> (visited on 11/06/2026).
- [31] Chris Olah, Nick Cammarata, Ludwig Schubert, Gabriel Goh, Michael Petrov and Shan Carter. ‘Zoom In: An Introduction to Circuits’. In: *Distill* 5:3 (Mar. 2020), e00024.001. ISSN: 2476-0757.

- DOI: [10.23915/distill.00024.001](https://doi.org/10.23915/distill.00024.001). URL: <https://distill.pub/2020/circuits/zoom-in> (visited on 04/05/2026).
- [32] Bruno A. Olshausen and David J. Field. ‘Sparse Coding with an Overcomplete Basis Set: A Strategy Employed by V1?’ In: *Vision Research* 37.23 (Dec. 1997), pp. 3311–3325. ISSN: 0042-6989. DOI: [10.1016/S0042-6989\(97\)00169-7](https://doi.org/10.1016/S0042-6989(97)00169-7). URL: <https://www.sciencedirect.com/science/article/pii/S0042698997001697> (visited on 28/05/2026).
- [33] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei and Ilya Sutskever. *Language Models Are Unsupervised Multitask Learners*. Tech. rep. OpenAI, 2019. (Visited on 13/06/2026).
- [34] Senthooran Rajamanoharan, Arthur Conmy, Lewis Smith, Tom Lieberum, Vikrant Varma, János Kramár, Rohin Shah and Neel Nanda. ‘Improving Sparse Decomposition of Language Model Activations with Gated Sparse Autoencoders’. In: *Advances in Neural Information Processing Systems* 37 (Dec. 2024), pp. 775–818. DOI: [10.52202/079017-0024](https://doi.org/10.52202/079017-0024). URL: https://proceedings.neurips.cc/paper_files/paper/2024/hash/01772a8b0420baec00c4d59fe2fbace6-Abstract-Conference.html (visited on 05/06/2026).
- [35] Senthooran Rajamanoharan, Tom Lieberum, Nicolas Sonnerat, Arthur Conmy, Vikrant Varma, János Kramár and Neel Nanda. *Jumping Ahead: Improving Reconstruction Fidelity with JumpReLU Sparse Autoencoders*. Comment: v2: new appendix H comparing kernel functions & bug-fixes to pseudo-code in Appendix J v3: further bug-fix to pseudo-code in Appendix J. Aug. 2024. DOI: [10.48550/arXiv.2407.14435](https://doi.org/10.48550/arXiv.2407.14435). arXiv: [2407.14435](https://arxiv.org/abs/2407.14435) [cs]. URL: <http://arxiv.org/abs/2407.14435> (visited on 23/02/2026).
- [36] Danilo Jimenez Rezende, Shakir Mohamed and Daan Wierstra. ‘Stochastic Backpropagation and Approximate Inference in Deep Generative Models’. In: *Proceedings of the 31st International Conference on Machine Learning*. PMLR, June 2014, pp. 1278–1286. URL: <https://proceedings.mlr.press/v32/rezende14.html> (visited on 10/06/2026).
- [37] Marco Tulio Ribeiro, Sameer Singh and Carlos Guestrin. ‘“Why Should I Trust You?”: Explaining the Predictions of Any Classifier’. In: *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. KDD ’16. New York, NY, USA: Association for Computing Machinery, Aug. 2016, pp. 1135–1144. ISBN: 978-1-4503-4232-2. DOI: [10.1145/2939672.2939778](https://doi.org/10.1145/2939672.2939778). URL: <https://dl.acm.org/doi/10.1145/2939672.2939778> (visited on 01/06/2026).

- [38] Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser and Bjorn Ommer. ‘High-Resolution Image Synthesis with Latent Diffusion Models’. In: *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. New Orleans, LA, USA: IEEE, June 2022, pp. 10674–10685. ISBN: 978-1-6654-6946-3. DOI: [10.1109/CVPR52688.2022.01042](https://doi.org/10.1109/CVPR52688.2022.01042). URL: <https://ieeexplore.ieee.org/document/9878449/> (visited on 10/06/2026).
- [39] Cynthia Rudin. ‘Stop Explaining Black Box Machine Learning Models for High Stakes Decisions and Use Interpretable Models Instead’. In: *Nature Machine Intelligence* 1.5 (May 2019), pp. 206–215. ISSN: 2522-5839. DOI: [10.1038/s42256-019-0048-x](https://doi.org/10.1038/s42256-019-0048-x). URL: <https://www.nature.com/articles/s42256-019-0048-x> (visited on 01/06/2026).
- [40] Mayu Sakurada and Takehisa Yairi. ‘Anomaly Detection Using Autoencoders with Nonlinear Dimensionality Reduction’. In: *Proceedings of the MLSDA 2014 2nd Workshop on Machine Learning for Sensory Data Analysis*. MLSDA’14. New York, NY, USA: Association for Computing Machinery, Dec. 2014, pp. 4–11. ISBN: 978-1-4503-3159-3. DOI: [10.1145/2689746.2689747](https://doi.org/10.1145/2689746.2689747). URL: <https://dl.acm.org/doi/10.1145/2689746.2689747> (visited on 05/06/2026).
- [41] Ramprasaath R. Selvaraju, Michael Cogswell, Abhishek Das, Ramakrishna Vedantam, Devi Parikh and Dhruv Batra. ‘Grad-CAM: Visual Explanations from Deep Networks via Gradient-Based Localization’. In: *2017 IEEE International Conference on Computer Vision (ICCV)*. Oct. 2017, pp. 618–626. DOI: [10.1109/ICCV.2017.74](https://doi.org/10.1109/ICCV.2017.74). URL: <https://ieeexplore.ieee.org/document/8237336> (visited on 01/06/2026).
- [42] Irhum Shafkat. *Intuitively Understanding Variational Autoencoders*. Feb. 2018. URL: <https://medium.com/data-science/intuitively-understanding-variational-autoencoders-1bfe67eb5daf> (visited on 14/05/2026).
- [43] Karen Simonyan, Andrea Vedaldi and Andrew Zisserman. *Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps*. Apr. 2014. DOI: [10.48550/arXiv.1312.6034](https://doi.org/10.48550/arXiv.1312.6034). arXiv: [1312.6034 \[cs.CV\]](https://arxiv.org/abs/1312.6034). URL: <http://arxiv.org/abs/1312.6034> (visited on 11/06/2026).
- [44] Chung-En Sun, Tuomas Oikarinen, Berk Ustun and Tsui-Wei Weng. ‘Concept Bottleneck Large Language Models’. In: *International Conference on Learning Representations 2025* (May 2025), pp. 89371–89411. URL: https://proceedings.iclr.cc/paper_files/paper/2025/hash/de4ce91dfe56b919ee1c228d6a78f866-Abstract-Conference.html (visited on 01/06/2026).

- [45] Adly Templeton et al. *Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet*. URL: <https://transformer-circuits.pub/2024/scaling-monosemanticity/index.html#appendix-citation> (visited on 20/10/2025).
- [46] Robert Tibshirani. ‘Regression Shrinkage and Selection Via the Lasso’. In: *Journal of the Royal Statistical Society: Series B (Methodological)* 58.1 (Jan. 1996), pp. 267–288. ISSN: 0035-9246. DOI: [10.1111/j.2517-6161.1996.tb02080.x](https://doi.org/10.1111/j.2517-6161.1996.tb02080.x). URL: <https://doi.org/10.1111/j.2517-6161.1996.tb02080.x> (visited on 29/05/2026).
- [47] Francesco Tonolini, Bjørn Sand Jensen and Roderick Murray-Smith. ‘Variational Sparse Coding’. In: *Proceedings of The 35th Uncertainty in Artificial Intelligence Conference*. PMLR, Aug. 2020, pp. 690–700. URL: <https://proceedings.mlr.press/v115/tonolini20a.html> (visited on 26/03/2026).
- [48] Aaron van den Oord, Oriol Vinyals and Koray Kavukcuoglu. ‘Neural Discrete Representation Learning’. In: *Advances in Neural Information Processing Systems*. Vol. 30. Curran Associates, Inc., 2017. URL: <https://papers.nips.cc/paper/2017/hash/7a98af17e63a0ac09ce2e96d03992fbc-Abstract.html> (visited on 11/06/2026).
- [49] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser and Illia Polosukhin. ‘Attention Is All You Need’. In: *Advances in Neural Information Processing Systems*. Vol. 30. Curran Associates, Inc., 2017. URL: https://proceedings.neurips.cc/paper_files/paper/2017/hash/3f5ee243547dee91fbd053c1c4a845aa-Abstract.html (visited on 11/06/2026).
- [50] Pascal Vincent, Hugo Larochelle, Yoshua Bengio and Pierre-Antoine Manzagol. ‘Extracting and Composing Robust Features with Denoising Autoencoders’. In: *Proceedings of the 25th International Conference on Machine Learning - ICML ’08*. Helsinki, Finland: ACM Press, 2008, pp. 1096–1103. ISBN: 978-1-60558-205-4. DOI: [10.1145/1390156.1390294](https://doi.org/10.1145/1390156.1390294). URL: <http://portal.acm.org/citation.cfm?doid=1390156.1390294> (visited on 05/06/2026).
- [51] Sandra Wachter, Brent Mittelstadt and Chris Russell. *Counterfactual Explanations Without Opening the Black Box: Automated Decisions and the GDPR*. SSRN Scholarly Paper. Rochester, NY, Oct. 2017. DOI: [10.2139/ssrn.3063289](https://doi.org/10.2139/ssrn.3063289). Social Science Research Network: [3063289](https://papers.ssrn.com/abstract=3063289). URL: <https://papers.ssrn.com/abstract=3063289> (visited on 11/06/2026).
- [52] Lilian Weng. *From Autoencoder to Beta-VAE*. Aug. 2018. URL: <https://lilianweng.github.io/posts/2018-08-12-vae/> (visited on 05/06/2026).

- [53] Benjamin Wright and Lee Sharkey. ‘Addressing Feature Suppression in SAEs — AI Alignment Forum’. In: (Feb. 2024). URL: <https://www.alignmentforum.org/posts/3JuSjTZyMzaSeTxKk/addressing-feature-suppression-in-saes> (visited on 29/05/2026).

COLOPHON

This document was typeset using the typographical look-and-feel `classicthesis` developed by André Miede and Ivo Pletikosić. The style was inspired by Robert Bringhurst’s seminal book on typography “*The Elements of Typographic Style*”. `classicthesis` is available for both $\text{\LaTeX}$ and $\text{\LyX}$:

<https://bitbucket.org/amiede/classicthesis/>

Happy users of `classicthesis` usually send a real postcard to the author, a collection of postcards received so far is featured here:

<http://postcards.miede.de/>

Thank you very much for your feedback and contribution.